

CSCE 4263/5183  
Advanced Data Structures

## Lecture 15

### *Hashing Topics (continue)*

Fall 2025

Prof. Khoa Luu  
[khoaluu@uark.edu](mailto:khoaluu@uark.edu)

# Notice

- HW 4 – Submission deadline: Nov. 11, 2025
- Final Project – Q&A

# Outline

Previously, we considered means for calculating 32-bit hash values


- Explicitly defined hash values
- Implicitly calculated hash values

Practically, we will require a hash value on the range  $0, \dots, M - 1$ :

- The modulus operator %
- Review of bitwise operations
- The multiplicative method

## 9.3.1

# The hash process



# Properties

Necessary properties of this mapping function  $h_M$  are:

- 2a. Must be fast:  $\Theta(1)$
- 2b. The hash value must be *deterministic*
  - Given  $n$  and  $M$ ,  $h_M(n)$  must always return the same value
- 2c. If two objects are randomly chosen, there should be only a one-in- $M$  chance that they have the same value from 0 to  $M - 1$

# Modulus operator

Easiest method: return the value modulus  $M$

```
unsigned int hash_M( unsigned int n, unsigned int M ) {  
    return n % M;  
}
```

Unfortunately, calculating the modulus (or remainder) is expensive

- We can simplify this if  $M = 2^m$
- We can use logic operations
  - We will review bitwise operations: left and right shift and bit-wise and

# The bitwise operators: & << >>

Suppose I want to calculate

$$7985325 \% 100$$

The modulo is a power of ten:  $100 = 10^2$

- In this case, take the last **two** decimal digits: 25

Similarly,  $7985325 \% 10^3 = 325$

- We set the appropriate digits to 0:

$$\text{0000025 and } \text{0000325}$$

# The bitwise operators: & << >>

The same works in base 2:

$$100011100101_2 \% 10000_2$$

The modulo is a power of 2:  $10000_2 = 2^4$

- In this case, take the last **four** bits: 0101

Similarly,  $100011100101_2 \% 1000000_2 == 100101$ ,

- We set the appropriate digits to 0:

$$\textcolor{red}{00000000}101 \text{ and } \textcolor{blue}{000000}100101$$

## The bitwise operators: & << >>

To zero all but the last  $n$  bits, select the last  $n$  bits using *bitwise and*:

$1000\ 1110\ \textcolor{red}{0101}_2 \ \& \ 0000\ 0000\ \textcolor{red}{1111}_2 \rightarrow 0000\ 0000\ \textcolor{red}{0101}_2$

$1000\ 11\textcolor{blue}{10}\ 0101_2 \ \& \ 0000\ 00\textcolor{blue}{11}\ 1111_2 \rightarrow 0000\ 00\textcolor{blue}{10}\ 0101_2$

## The bitwise operators: & << >>

Similarly, multiplying or dividing by powers of 10 is easy:

7985325 \* 100

The multiplier is a power of ten:  $100 = 10^2$

- In this case, add **two** zeros: 798532500

Similarly,  $7985325 / 10^3 = 7985$

- Just add the appropriate number of zeros or remove the appropriate number of digits

## The bitwise operators: & << >>

The same works in base 2:

$$100011100101_2 * 10000_2$$

The modulo is a power of 2:  $10000_2 = 2^4$

- In this case, add **four** zeros:  $1000111001010000$

Similarly,  $100011100101_2 / 1000000_2 == 100011$

## The bitwise operators: & << >>

This can be done mechanically by shifting the bits appropriately:

$$\begin{aligned}1000\ 1110\ 0101_2 \ll 4 &== 1000\ 1110\ 0101\ \textcolor{red}{0000}_2 \\1000\ 11\textcolor{blue}{10}\ 0101_2 \gg 6 &== 10\ 0011_2\end{aligned}$$

Powers of 2 are now easy to calculate:

$$\begin{aligned}1_2 \ll 4 &== 1\ 0000_2 && // 2^4 = 16 \\1_2 \ll 6 &== 100\ 0000_2 && // 2^6 = 64\end{aligned}$$

## Modulo a power of two

The implementation using the modulus/remainder operator:

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    return n & ((1 << m) - 1);  
}
```

# Modulo a power of two

Problem:

- Suppose that the hash function  $h$  is always even
- An even number modulo a power of two is still even

Example: memory allocations are multiples of word size

- On a 64-bit computer, addresses returned by new will be multiples of 8
- Thus, if  $x \neq y$ ,  $h(x) \neq h(y)$
- However, the probability that  $h_M(h(x)) = h_M(h(y))$  is one in  $M/8$ 
  - This is not one in  $M$

## Modulo a power of two

For some objects, it is worse:

- Instance of `Single_node<int>` on linux always have `10000` as the last five bits
  - This increases the probability of a collision to one in  $M/32$

Fortunately, the multiplicative method resolves this issue

# Multiplication techniques

We need to obfuscate the bits

- The most common method to obfuscate bits is multiplication
- Consider how one bit can affect an entire range of numbers in the result:

$$\begin{array}{r} 10100\textcolor{red}{1}11 \\ \times \underline{11010011} \\ 10100\textcolor{red}{1}11 \end{array}$$

10100~~1~~11

C++ truncates the product of two  $n$ -bit numbers to  $n$  bits

$$\begin{array}{r} 10100\textcolor{red}{1}11 \\ 10100\textcolor{red}{1}11 \\ + \underline{10100\textcolor{red}{1}11} \\ 100010\textcolor{red}{1}110100101 \end{array}$$

The *avalanche* effect: when changing one bits has the potential of affecting all bits in the result:

$$\begin{aligned} 10100011 \times 11010011 \\ = 1000011001011001 \end{aligned}$$

# Multiplication techniques

Multiplication can be used in encryption as with IDEA:

```
unsigned short a, b, c, d; // 64 bit number broken into four 16 bits
unsigned short z[9][6];    // 52 numbers based on key (2 unused)

for ( int i = 0; i < 8; ++i ) {
    a *= z[i][0];  b += z[i][1];  c += z[i][2];  d *= z[i][3];
    e = z[i][4]*(a^c);
    f = ( e + (b^d) )*z[i][5];
    e += f;
    a ^= f;        tmp = b;      b = c^f;      c = tmp^e;      d ^= e;
}

a *= z[8][0];  b += z[8][1];  c += z[8][2];  d *= z[8][3];
```

# The multiplicative method

Multiplying by a fixed constant is a reasonable method

- Take the middle  $m$  bits of  $Cn$ :

```
unsigned int const C = 581869333; // some number

unsigned int hash_M( unsigned int n, unsigned int m ) {
    unsigned int shift = (32 - m)/2;
    return ((C*n) >> shift) & ((1 << m) - 1);
}
```

# The multiplicative method

Suppose that the value  $m = 10$  ( $M = 1024$ ) and  $n = 42$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}
```

# The multiplicative method

$$m = 10$$

$$n = 42$$

First calculate the shift

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}  
shift = 11
```

# The multiplicative method

$$m = 10$$

$$n = 42$$

Next,  $n = 42$  or  $101010_2$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}  
shift = 11
```

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	1	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

# The multiplicative method

$m = 10$

$n = 42$

Calculate  $Cn$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {
    unsigned int shift = (32 - m)/2;
    return ((C*n) >> shift) & ((1 << m) - 1);
}
shift = 11
```

1	0	1	1	0	0	0	0	1	0	1	0	0	1	1	0	0	0	0	1	1	0	0	1	0	1	1	1	0	0	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


# The multiplicative method

 $m = 10$ 
 $n = 42$ 

Right shift this value 11 bits—equivalent to dividing by  $2^{11}$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {
    unsigned int shift = (32 - m)/2;
    return ((C*n) >> shift) & ((1 << m) - 1);
}
shift = 11
```



# The multiplicative method

$$\begin{aligned}m &= 10 \\n &= 42\end{aligned}$$

## Next, start with 1

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}
```

0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 1 0 1 0 0 1 1 1 0 0 0 0 0 1 1


# The multiplicative method

 $m = 10$  $n = 42$ 

Left shift 1  $m = 10$  bits yielding  $2^{10}$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}
```



# The multiplicative method

$$m = 10$$

$n = 42$

Subtracting 1 yields  $m = 10$  ones

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}
```

0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 1 0 1 0 0 1 1 0 0 0 0 1 1

# The multiplicative method

$$m = 10$$

$$n = 42$$

Taken the bitwise to clear all but the last 10 bits

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {
    unsigned int shift = (32 - m)/2;
    return ((C*n) >> shift) & ((1 << m) - 1);
}
```




# The multiplicative method

We have extracted the middle  $m = 10$  bits—a number in  $0, \dots, 1023$

```
const unsigned int C = 581869333; // some number
```

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    unsigned int shift = (32 - m)/2;  
    return ((C*n) >> shift) & ((1 << m) - 1);  
}
```



$$h_M(42) = 195$$

# Dealing with signed integers

In some cases (Java) unsigned integers are either not available

Our function would be implemented as follows:

```
int hash_M( int n, int M ) {  
    return n % M;  
}
```

Is this sufficient to yield a value on  $0, \dots, M - 1$ ?

# Dealing with signed integers

The modulus operator  $\%$  satisfies the following equality:

$$n == M*(n/M) + n\%M$$

The division operator  $n/M$  truncates any fractional part

- It rounds towards zero

For positive  $n$  and  $M$ , this yields positive values:

- If  $n = 13$  and  $M = 8$ ,  $13/8 \rightarrow 1$  and  $8*1 \rightarrow 8$   
 $\therefore 13 \% 8 \rightarrow 5$

For a negative value of  $n$ :

- If  $n = -11$  and  $M = 8$ ,  $-11/8 \rightarrow -1$  and  $8*(-1) = -8$   
 $\therefore (-11) \% 8 \rightarrow -3$

# Dealing with signed integers


Our mapping function could be implemented as follows:

```
int hash_M( int n, int M ) {  
    int hash_value = n % M;  
    return (hash_value >= 0) ? hash_value : hash_value + M;  
}
```


# Dealing with signed integers

Why  $+ M$  and not `std::abs( M )`?

- First, adding results in periodicity:



- Using the absolute value is not periodic



# Dealing with signed integers

Also, not all integers have absolute values:

```
int main() {
    int n1 = -2147483648; // -2^31
    int n2 = -2147483647; // -(2^31 - 1)

    cout << std::abs( n1 ) << endl;
    cout << std::abs( n2 ) << endl;

    return 0;
}
```

The output is:

```
-2147483648
2147483647
```

# Summary

This topic covered next step in creating a hash value:

- Taking a 32-bit number and mapping it down to  $0, \dots, M - 1$
- One can use the modulus; however, using the middle  $m$  bits of the number times a large prime avoids some weaknesses of the modulus

We must now deal with *collisions*: the reality that two objects may hash to the same value

# **Chained Hash Tables**

We have:

- Discussed techniques for hashing
- Discussed mapping down to a given range  $0, \dots, M - 1$

Now we must deal with collisions

- Numerous techniques exist
- Containers in general
  - Specifically linked lists

# Background

First, a review:

- We want to store objects in an array of size  $M$
- We want to quickly calculate the bin where we want to store the object
  - We came up with hash functions—hopefully  $\Theta(1)$
  - Perfect hash functions (no collisions) are difficult to design
- We will look at some schemes for dealing with collisions

# Implementation

Consider associating each bin with a linked list:

```
template <class Type>
class Chained_hash_table {
    private:
        int table_capacity;
        int table_size;
        Single_list<Type> *table;


        unsigned int hash( Type const & );

    public:
        Chained_hash_table( int = 16 );
        int count( Type const & ) const;
        void insert( Type const & );
        int erase( Type const & );
        // ...
};
```

# Implementation

The constructor creates an array of  $n$  linked lists

```
template <class Type>
Chained_hash_table::Chained_hash_table( int n ) :
    table_capacity( std::max( n, 1 ) ),
    table_size( 0 ),
    table( new Single_list<Type>[table_capacity] ) {
    // empty constructor
}
```




# Implementation

The function `hash` will determine the bin of an object:

```
template <class Type>
int Chained_hash_table::hash( Type const &obj ) {
    return hash_M( obj.hash(), capacity() );
}
```

Recall:

- `obj.hash()` returns a 32-bit non-negative integer
- `unsigned int hash_M( obj, M )` returns a value in  $0, \dots, M - 1$




# Implementation

Other functions mapped to corresponding linked list functions:


```
template <class Type>
void Chained_hash_table::insert( Type const &obj ) {
    unsigned int bin = hash( obj );

    if ( table[bin].count( obj ) == 0 ) {
        table[bin].push_front( obj );
        ++table_size;
    }
}
```



# Implementation


```
template <class Type>
int Chained_hash_table::count( Type const &obj ) const {
    return table[hash( obj )].count( obj );
}
```



# Implementation

```
template <class Type>
int Chained_hash_table::erase( Type const &obj ) {
    unsigned int bin = hash( obj );


    if ( table[bin].erase( obj ) ) {
        --table_size;
        return 1;
    } else {
        return 0;
    }
}
```



# Example

As an example, let's store hostnames and allow a fast look-up of the corresponding IP address


- We will choose the bin based on the host name
- Associated with the name will be the IP address
- *E.g.*, ("optimal", 129.97.94.57)



# Example


We will store strings and the hash value of a string will be the last 3 bits of the first character in the host name

- The hash of "optimal" is based on "o"



# Example

The following is a list of the binary representation of each letter:  
– "a" is 1 and it cycles from there...




a	01100 <b>001</b>	n	01101 <b>110</b>
b	01100 <b>010</b>	o	01101 <b>111</b>
c	01100 <b>011</b>	p	01110 <b>000</b>
d	01100 <b>100</b>	q	01110 <b>001</b>
e	01100 <b>101</b>	r	01110 <b>010</b>
f	01100 <b>110</b>	s	01110 <b>011</b>
g	01100 <b>111</b>	t	01110 <b>100</b>
h	01101 <b>000</b>	u	01110 <b>101</b>
i	01101 <b>001</b>	v	01110 <b>110</b>
j	01101 <b>010</b>	w	01110 <b>111</b>
k	01101 <b>011</b>	x	01111 <b>000</b>
l	01101 <b>100</b>	y	01111 <b>001</b>
m	01101 <b>101</b>	z	01111 <b>010</b>

# Example


Our hash function is

```
unsigned int hash( string const &str ) {  
    // the empty string "" is hashed to 0  
    if str.length() == 0 ) {  
        return 0;  
    }  
    return str[0] & 7;  
}
```




# Example

Starting with an array of 8 empty linked lists



# Example


The pair ("optimal", 129.97.94.57) is entered into bin  
 $01101\textcolor{red}{111} = 7$



# Example

Similarly, as "c" hashes to 3


- The pair ("cheetah", 129.97.94.45) is entered into bin 3



# Example

The "w" in Wellington also hashes to 7

- ("wellington", 129.97.94.42) is entered into bin 7



# Example


Why did I use `push_front` from the linked list?

- Do I have a choice?
- A good heuristic is

“unless you know otherwise, data which has been

accessed recently will be accessed again in the near future”


- It is easiest to access data at the front of a linked list



Heuristics include rules of thumb,  
educated guesses, and intuition


# Example

Similarly we can insert the host names "augustin" and "lowpower"




# Example

If we now wanted the IP address for "optimal", we would simply hash "optimal" to 7, walk through the linked list, and access 129.97.94.57 when we access the node containing the relevant string




# Example

Similarly, "ashok" and "vlach" are entered into bin 7



# Example


Inserting "ims", "jab", and "cad" doesn't even out the bins



# Example

Indeed, after 21 insertions, the linked lists are becoming rather long

- We were looking for  $\Theta(1)$  access time, but accessing something in a linked list with  $k$  objects is  $O(k)$



## Load Factor

To describe the length of the linked lists, we define the *load factor* of the hash table:

$$\lambda = \frac{n}{M}$$

This is the average number of objects per bin

- This assumes an even distribution

Right now, the load factor is  $\lambda = 21/8 = 2.625$

- The average bin has 2.625 objects

## Load Factor

If the load factor becomes too large, access times will start to increase:  $\mathbf{O}(\lambda)$


The most obvious solution is to double the size of the hash table

- Unfortunately, the hash function must now change
- In our example, the doubling the hash table size requires us to take, for example, the last four bits

# Doubling Size


The load factor is now  $\lambda = 1.3125$

- Unfortunately, the distribution hasn't improved much



# Doubling Size

There is significant *clustering* in bins 2 and 3 due to the choice of host names



# Choosing a Good Hash Function

We choose a very poor hash function:

- We looked at the first letter of the host name

Unfortunately, all these are also actual host names:

ultra7 ultra8 ultra9 ultra10 ultra11  
ultra12 ultra13 ultra14 ultra15 ultra16 ultra17  
blade1 blade2 blade3 blade4 blade5

This will cause clustering in bins 2 and 5

- Any hash function based on anything other than every letter will cause clustering


# Choosing a Good Hash Function

Let's go back to the hash function defined previously:

```
unsigned int hash( string const &str ) {  
    unsigned int hash_value = 0;  
  
    for ( int k = 0; k < str.length(); ++k ) {  
        hash_value = 12347*hash_value + str[k];  
    }  
  
    return hash_value;  
}
```

# Choosing a Good Hash Function


This hash function yields a much nicer distribution:



# Choosing a Good Hash Function



When we insert the names of all 55 names, we would have a load factor  $\lambda = 3.4375$

- Clearly there are not exactly 3.4375 objects per bin
- How can we tell if this is a good hash function?
- Can we expect exactly 3.4375 objects per bin?
  - Clearly no...
- The answer is with statistics...



# Choosing a Good Hash Function

We would expect the number of bins which hold  $k$  objects to approximately follow a Poisson distribution



# Problems with Linked Lists

One significant issue with chained hash tables using linked lists

- It requires extra memory
- It uses dynamic memory allocation

Total memory:

16 bytes

- A pointer to an array, initial and current number of bins, and the size
- +  $12M$  bytes ( $8M$  if we remove count from `Single_list`)
- +  $8n$  bytes if each object is 4 bytes

## Problems with linked lists

For faster access, we could replace each linked list with an AVL tree (assuming we can order the objects)

- The access time drops to  $\mathbf{O}(\ln(\lambda))$
- The memory requirements are increased by  $\Theta(n)$ , as each node will require two pointers

We could look at other techniques:

- Scatter tables: use other bins and link the bins
- Use an alternate memory allocation model

# Black Board Example

Use the hash function

```
unsigned int hash( unsigned int n ) { return n % 10; }
```

to enter the following 15 numbers into a hash table with 10 bins:

534, 415, 465, 459, 869, 442, 840, 180, 450, 265, 23, 946, 657, 3, 29

# Summary

The easiest way to deal with collisions is to associate each bin with a container

We looked at bins of linked lists

- The example used host names and IP addresses
- We defined the load factor  $\lambda = n/M$
- Discussed doubling the number of bins
- Our goals are to choose a good hash function and to keep the load factor low
- We discussed alternatives

Next we will see a different technique using only one array of bins: open addressing