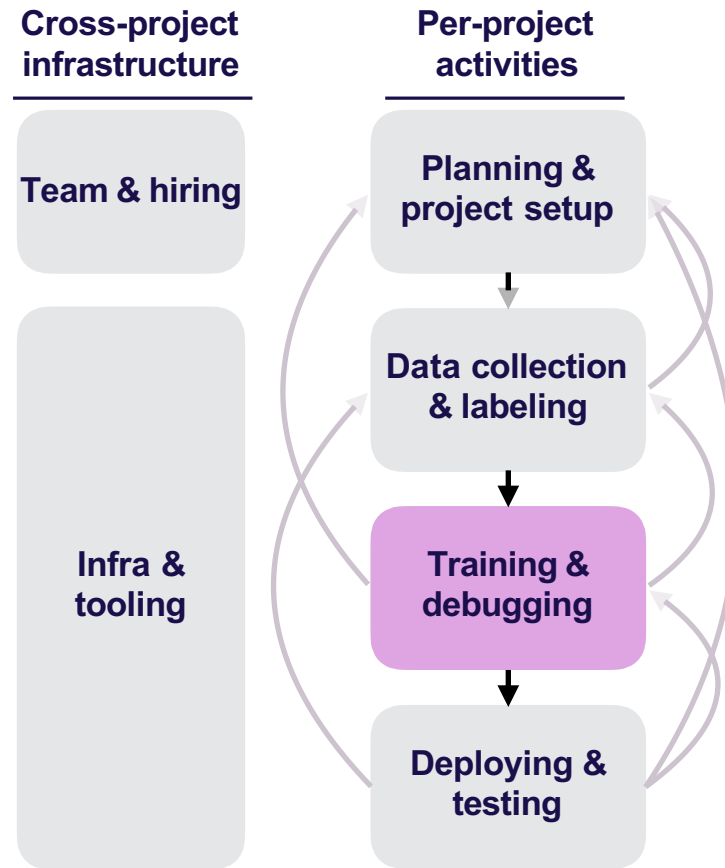


# Troubleshooting Deep Neural Networks

# Class Annoucement

- HW1 Grading: Released
- HW2 Submission: due Feb 14 (tomorrow)
- HW3: Release next week
- Midterm and Final Exams: Team Projects
- **Invited Talk: Walmart (Thu, Feb. 20)**

# Lifecycle of a ML project



# Why talk about DL troubleshooting?



XKCD, <https://xkcd.com/1838/>

# Why talk about DL troubleshooting?


**Common sentiment among practitioners:**

**80-90%** of time debugging and tuning

**10-20%** deriving math or implementing things

# Why is DL troubleshooting so hard?



# Suppose you can't reproduce a result



He, Kaiming, et al. "Deep residual learning for image recognition."  
*Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016.

# Suppose you can't reproduce a result


## Your learning curve



He, Kaiming, et al. "Deep residual learning for image recognition."  
*Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016.



# Suppose you can't reproduce a result




(a)


[https://papers.neurips.cc/paper\\_files/paper/2022/file/dfc310e81992d2e4cedc09ac47eff13e-Paper-Conference.pdf](https://papers.neurips.cc/paper_files/paper/2022/file/dfc310e81992d2e4cedc09ac47eff13e-Paper-Conference.pdf)

# Suppose you can't reproduce a result

## Grokking



(a)



(b)

Figure 8: Left: Training curves for a run on MNIST, in the setting where we observe grokking. Right: Phase diagram with the four phases of learning dynamics on MNIST.

# Suppose you can't reproduce a result

---

## Towards Understanding Grokking: An Effective Theory of Representation Learning

---

Ziming Liu, Ouail Kitouni, Niklas Nolte, Eric J. Michaud, Max Tegmark, Mike Williams  
Department of Physics, Institute for AI and Fundamental Interactions, MIT  
{zmliu,kitouni,nnolte,ericjm,tegmark,mwill}@mit.edu


### Abstract

We aim to understand *grokking*, a phenomenon where models generalize long after overfitting their training set. We present both a *microscopic* analysis anchored by an effective theory and a *macroscopic* analysis of phase diagrams describing learning performance across hyperparameters. We find that generalization originates from structured representations whose training dynamics and dependence on training set size can be predicted by our effective theory in a toy setting. We observe empirically the presence of four learning phases: *comprehension*, *grokking*, *memorization*, and *confusion*. We find representation learning to occur only in a “Goldilocks zone” (including comprehension and grokking) between memorization and confusion. We find on transformers the grokking phase stays closer to the memorization phase (compared to the comprehension phase), leading to delayed generalization. The Goldilocks phase is reminiscent of “intelligence from starvation” in Darwinian evolution, where resource limitations drive discovery of more efficient solutions. This study not only provides intuitive explanations of the origin of grokking, but also highlights the usefulness of physics-inspired tools, e.g., effective theories and phase diagrams, for understanding deep learning.

# Why is your performance worse?


Poor model  
performance

# Why is your performance worse?



# Most DL bugs are invisible


```
1 features = glob.glob('path/to/features/*')  
2 labels = glob.glob('path/to/labels/*')  
3 train(features, labels)
```




# Most DL bugs are invisible

Labels out of order!

```
1 features = glob.glob('path/to/features/*')
2 labels = glob.glob('path/to/labels/*')
3 train(features, labels)
```




# Why is your performance worse?






# Why is your performance worse?




# Models are sensitive to hyperparameters




*Andrej Karpathy, CS231n course notes*

# Models are sensitive to hyperparameters




Andrej Karpathy, CS231n course notes

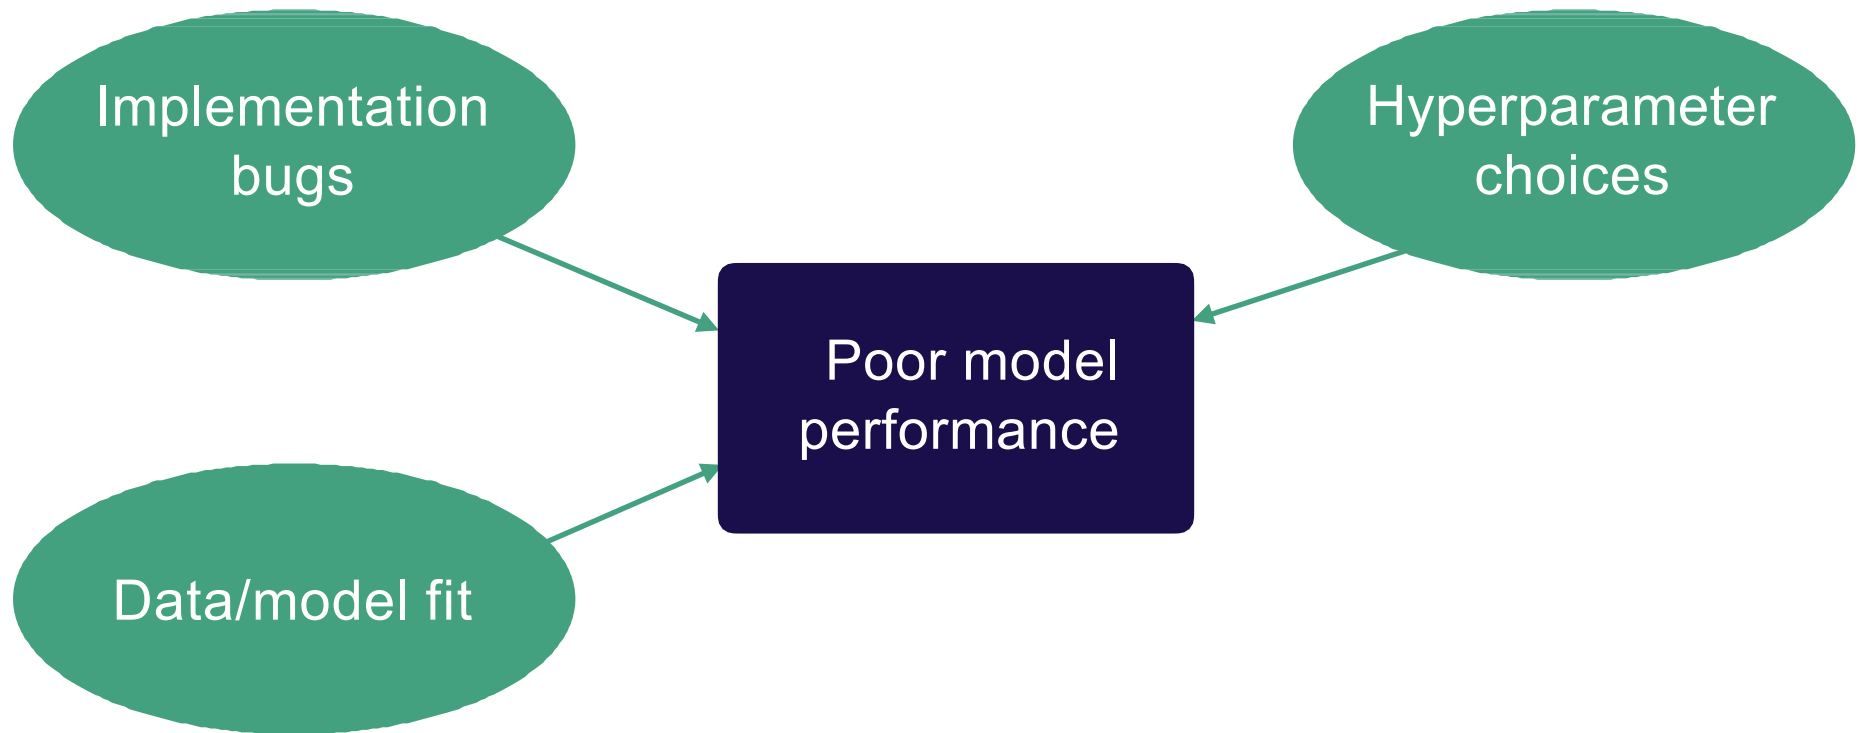


He, Kaiming, et al. "Delving deep into rectifiers: Surpassing human-level performance on imagenet classification." *Proceedings of the IEEE international conference on computer vision*. 2015.

# Why is your performance worse?




# Why is your performance worse?



# Data / model fit

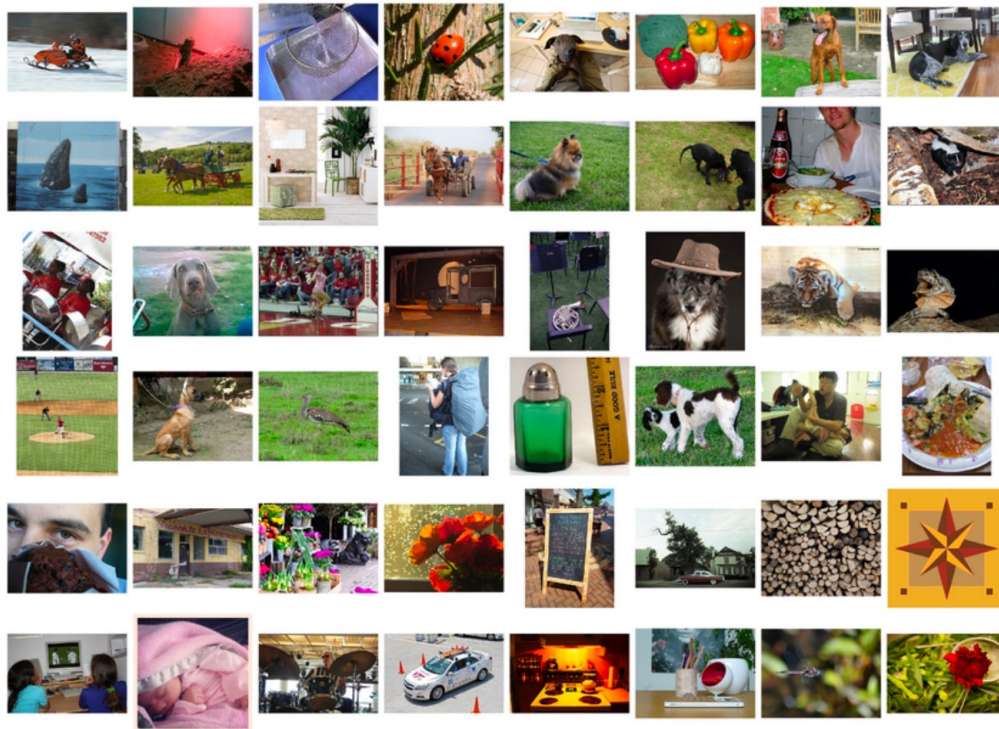
## Data from the paper: ImageNet





# Data / model fit

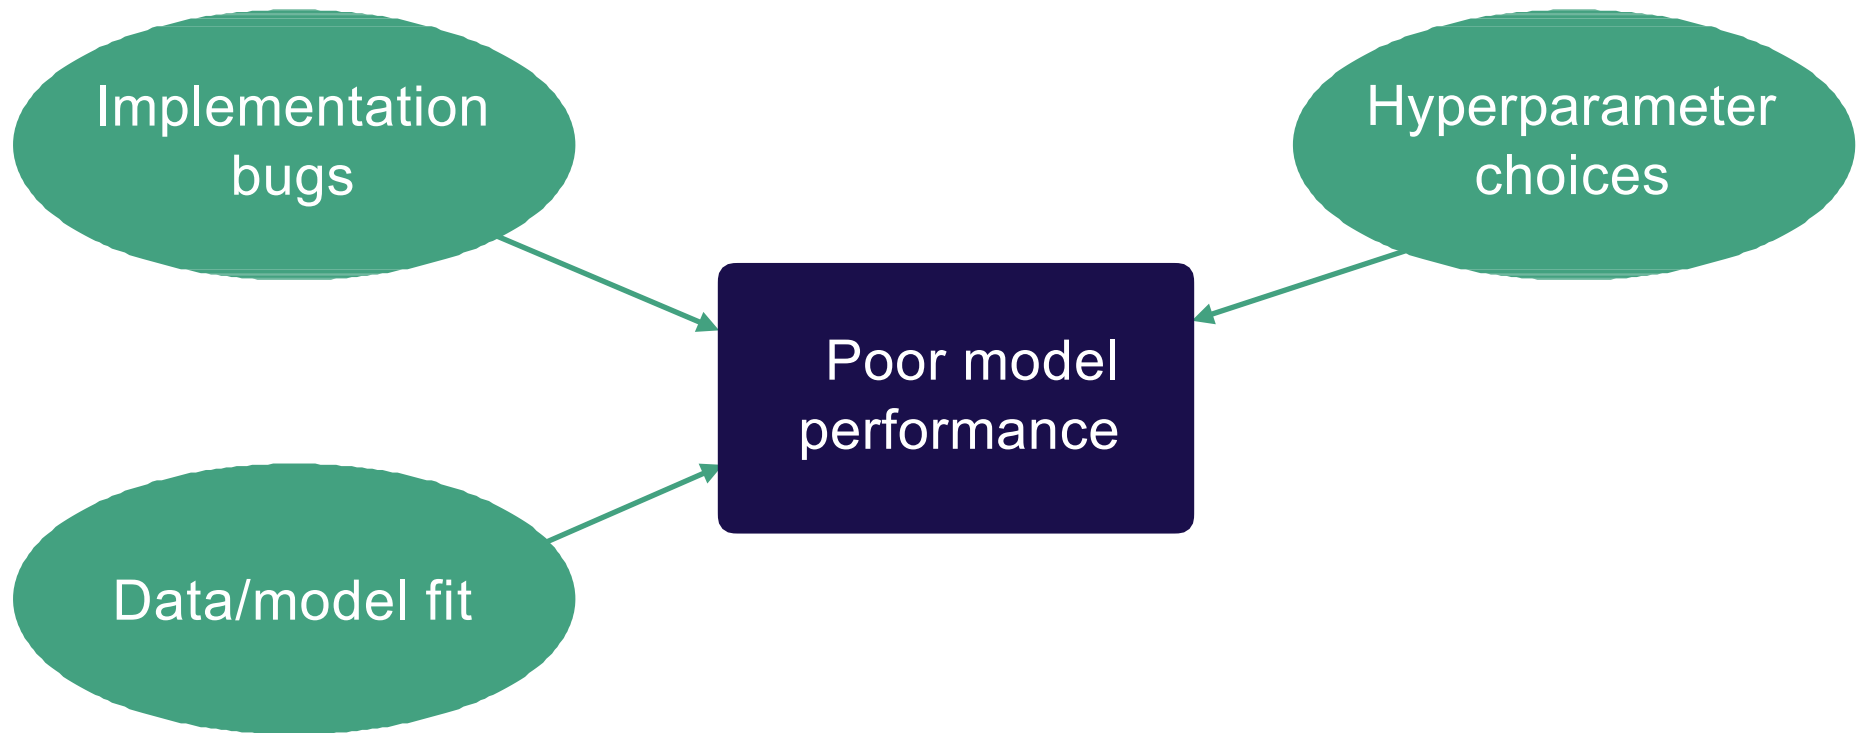
**Data from the paper: ImageNet**



**Yours: self-driving car images**




# Why is your performance worse?






# Why is your performance worse?



# Constructing good datasets is hard



Slide from Andrej Karpathy's talk "Building the Software 2.0 Stack" at TrainAI 2018, 5/10/2018

# Common dataset construction issues

- Not enough data
- Class imbalances
- Noisy labels
- Train/ test from different distributions
- etc

# Takeaways: why is troubleshooting hard?

- Hard to tell if you have a bug
- Lots of possible sources for the same degradation in performance
- Results can be sensitive to small changes in hyperparameters and dataset makeup

# Strategy for DL troubleshooting

# Key mindset for DL troubleshooting

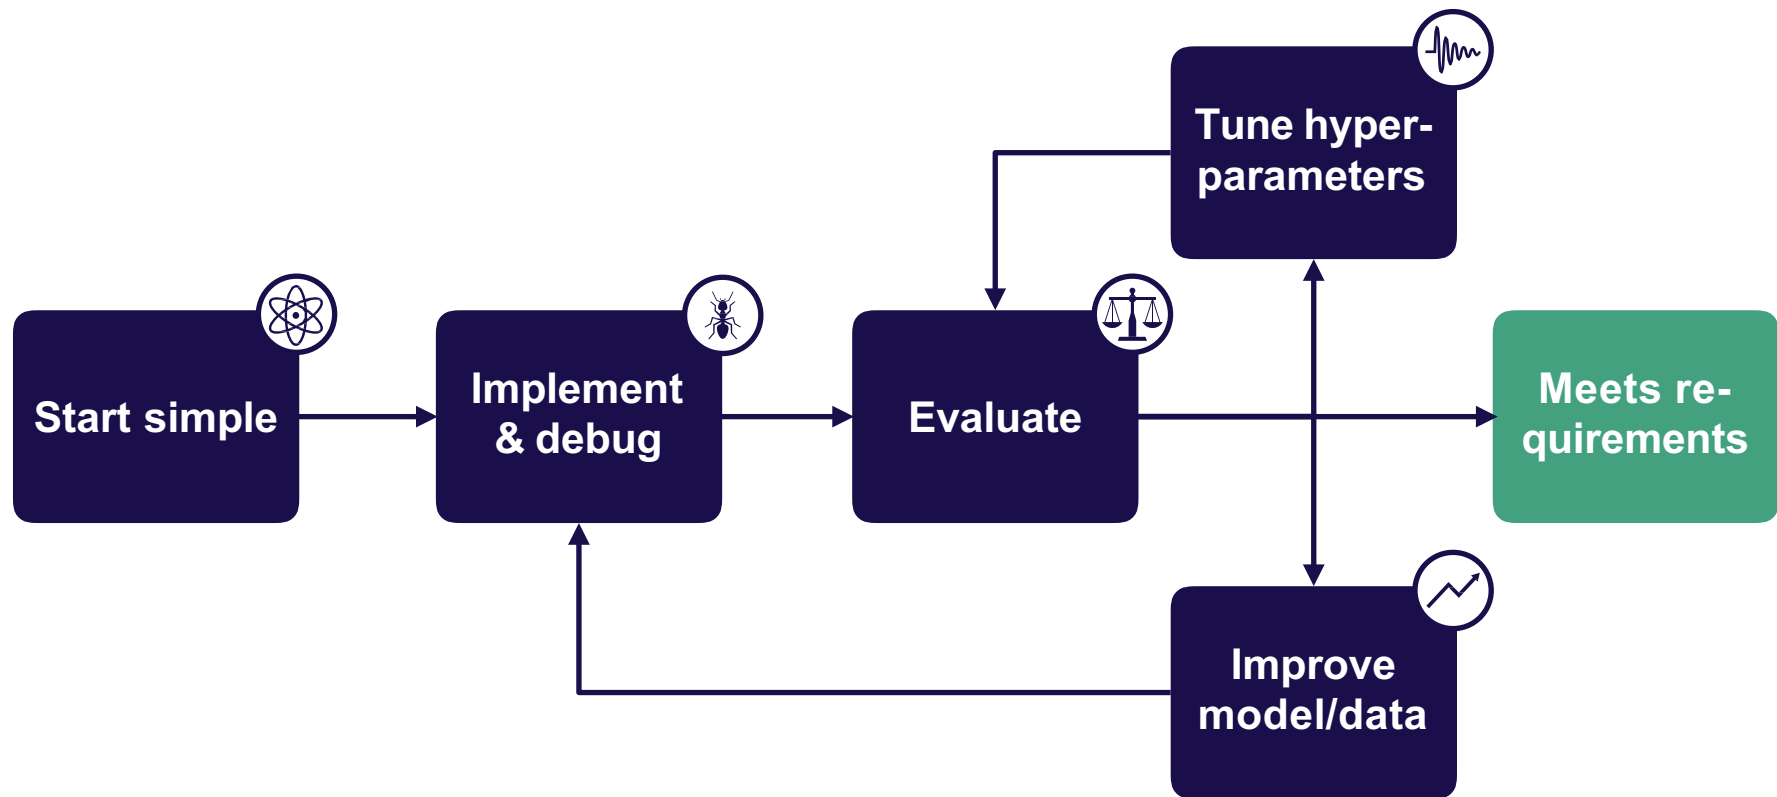
**Pessimism**

# Key idea of DL troubleshooting

**Since it's hard to  
disambiguate errors...**

**...Start simple and gradually  
ramp up complexity**

# Strategy for DL troubleshooting





# Quick summary



**Start  
simple**

- Choose the simplest model & data possible (e.g., LeNet on a subset of your data)

# Quick summary



**Start  
simple**



**Implement  
& debug**

- Choose the simplest model & data possible (e.g., LeNet on a subset of your data)
- Once model runs, overfit a single batch & reproduce a known result

# Quick summary



**Start  
simple**

- Choose the simplest model & data possible (e.g., LeNet on a subset of your data)



**Implement  
& debug**

- Once model runs, overfit a single batch & reproduce a known result



**Evaluate**

- Apply the bias-variance decomposition to decide what to do next

# Quick summary



**Start  
simple**

- Choose the simplest model & data possible (e.g., LeNet on a subset of your data)



**Implement  
& debug**

- Once model runs, overfit a single batch & reproduce a known result



**Evaluate**






- Apply the bias-variance decomposition to decide what to do next



**Tune hyp-  
eparams**

- Use coarse-to-fine random searches

# Quick summary

-  **Start simple**
  - Choose the simplest model & data possible (e.g., LeNet on a subset of your data)
-  **Implement & debug**
  - Once model runs, overfit a single batch & reproduce a known result
-  **Evaluate**
  - Apply the bias-variance decomposition to decide what to do next
-  **Tune hyp-params**
  - Use coarse-to-fine random searches
-  **Improve model/data**
  - Make your model bigger if you underfit; add data or regularize if you overfit

# We'll assume you already have...

- Initial test set
- A single metric to improve
- Target performance based on human-level performance, published results, previous baselines, etc

# We'll assume you already have...

## Running example

- Initial test set
- A single metric to improve
- Target performance based on human-level performance, published results, previous baselines, etc



0 (no pedestrian)

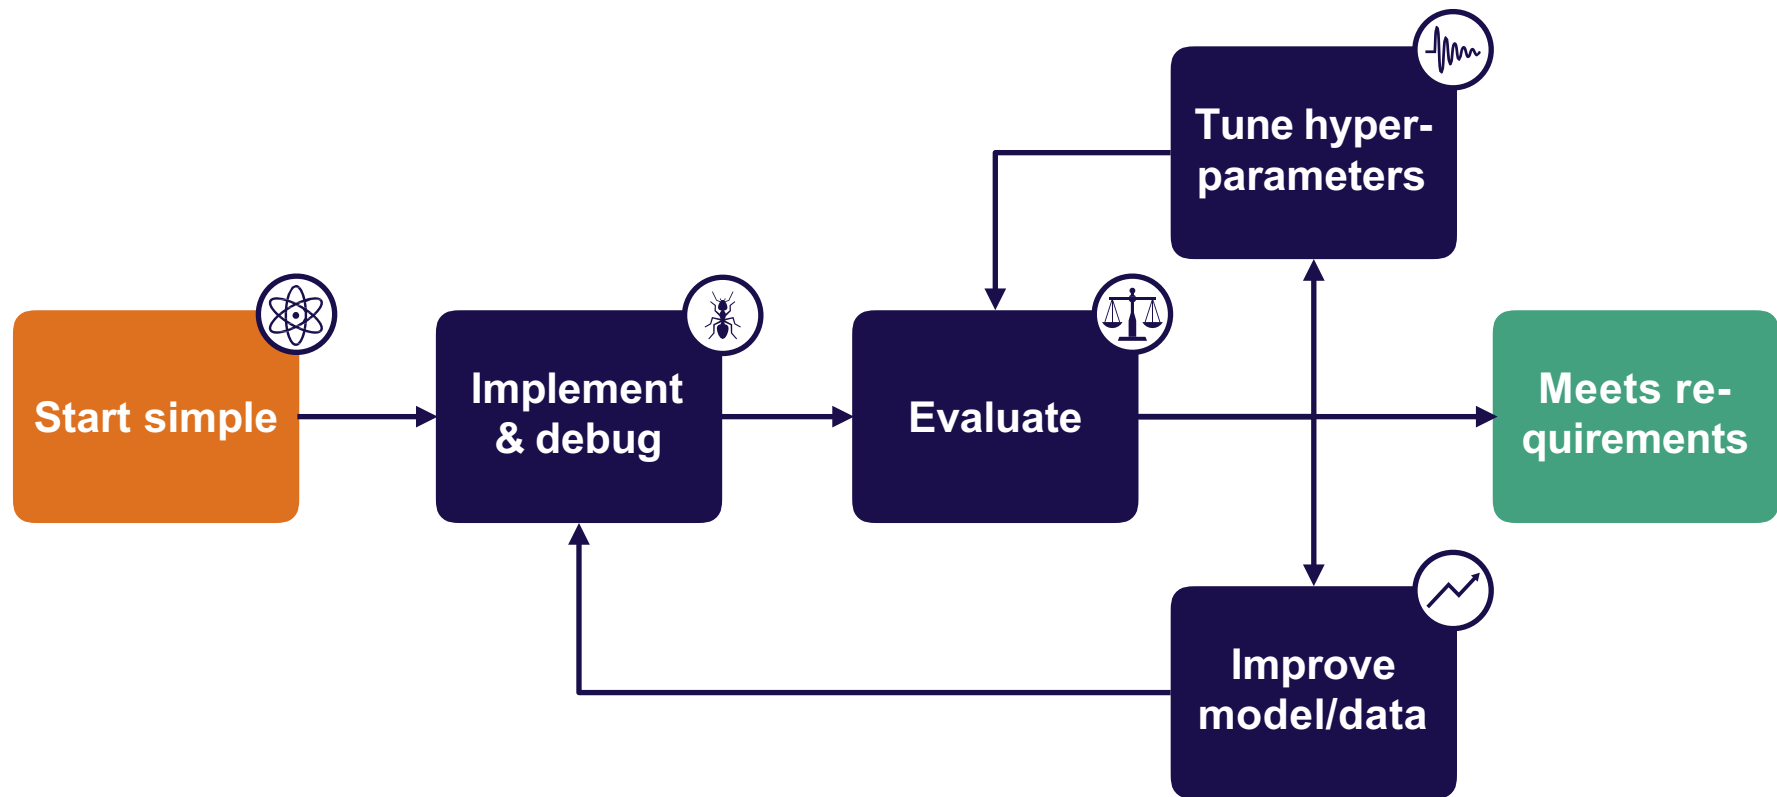
1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Questions?

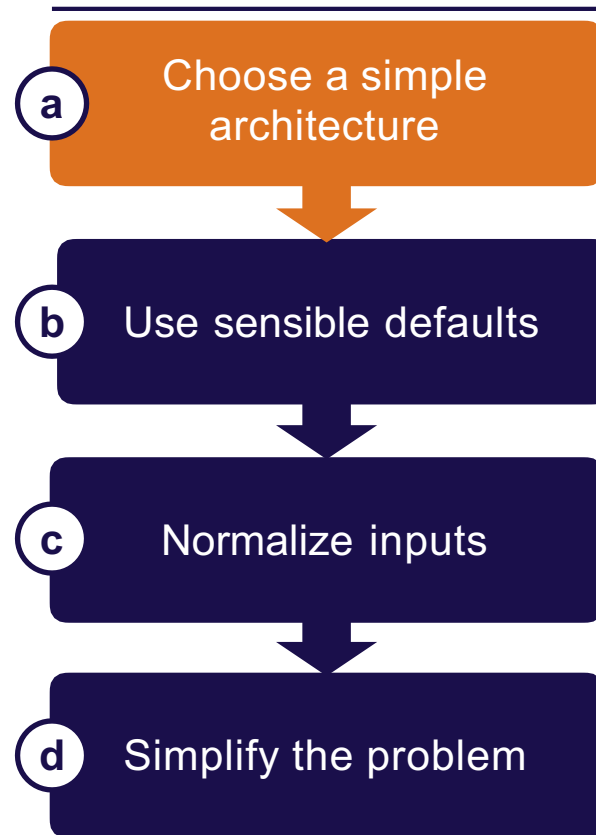


# Strategy for DL troubleshooting

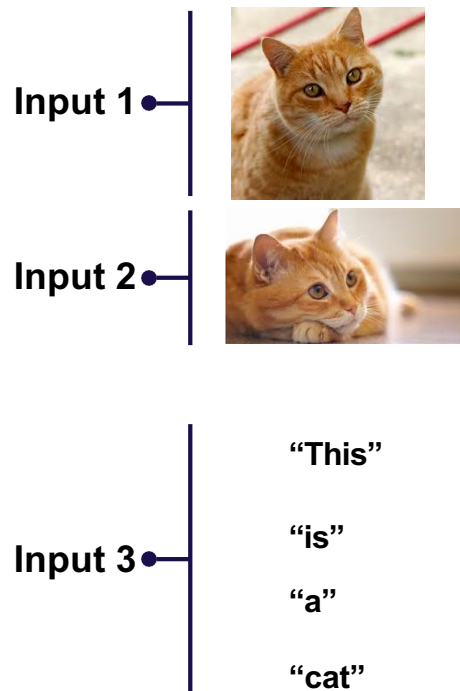


# Starting simple

## Steps




# Dealing with multiple input modalities



# Dealing with multiple input modalities

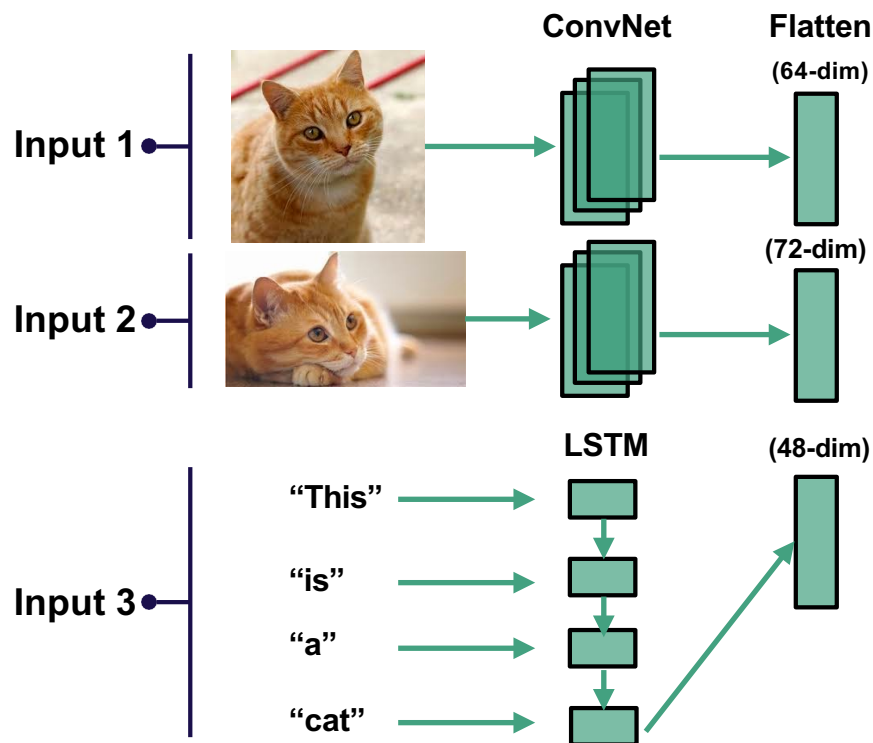
## 1. Map each into a lower dimensional feature space

---



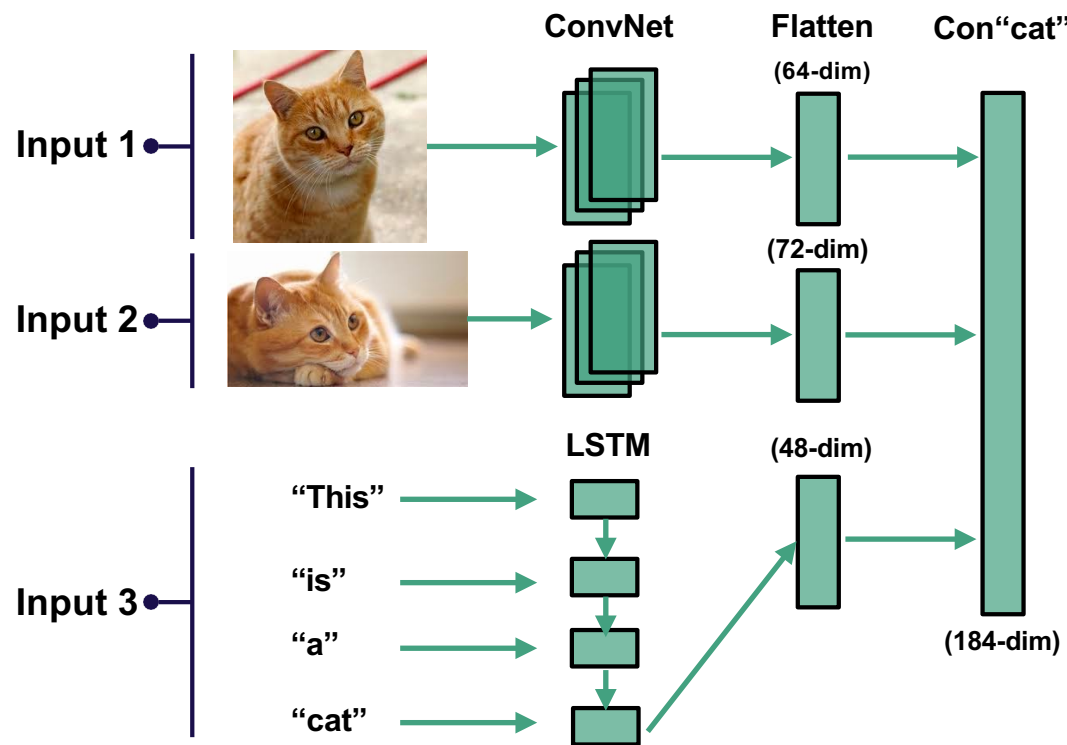
# Dealing with multiple input modalities

## 1. Map each into a lower dimensional feature space



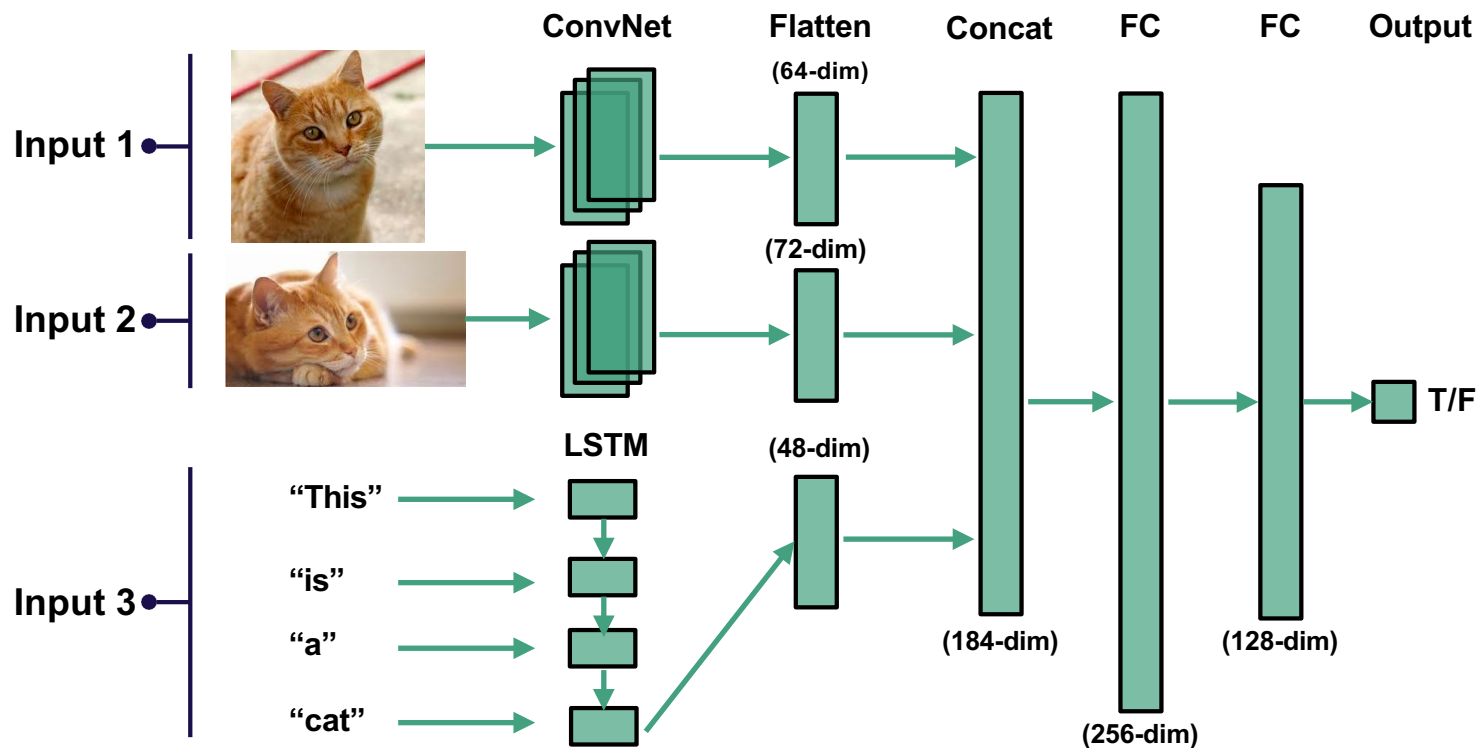
# Dealing with multiple input modalities

## 2. Concatenate



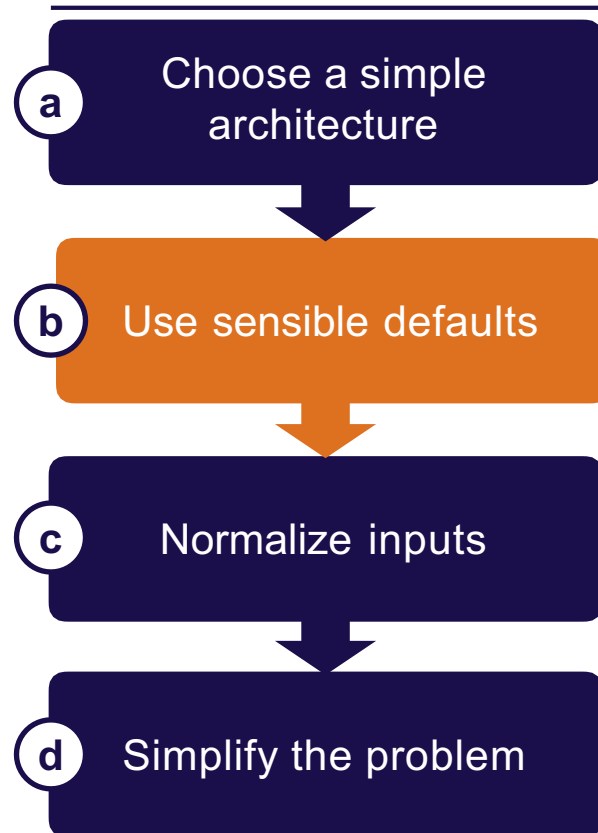
# Dealing with multiple input modalities

## 3. Pass through fully connected layers to output



# Starting simple

## Steps



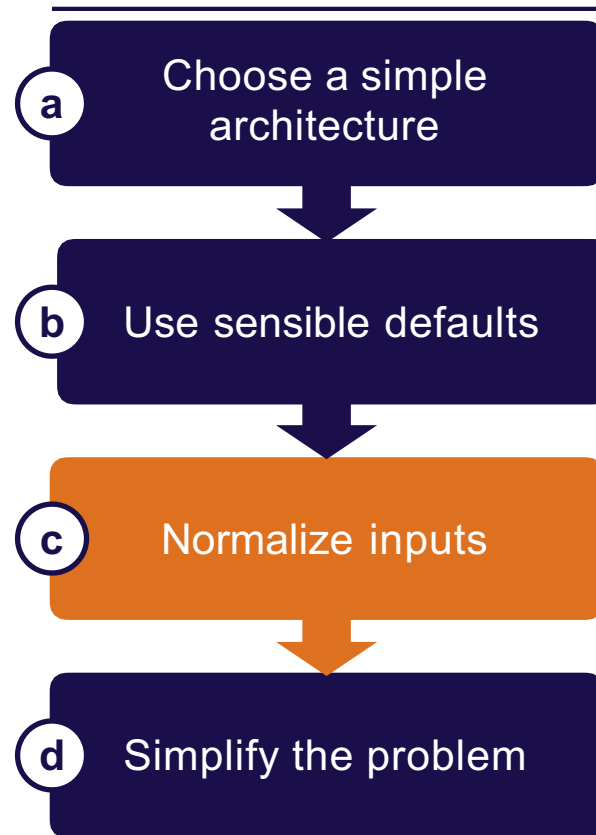


# Recommended network / optimizer defaults

- **Optimizer:** Adam optimizer with learning rate  $3e-4$
- **Activations:** relu (FC and Conv models), tanh (LSTMs)
- **Initialization:** He et al. normal (relu), Glorot normal (tanh)
- **Regularization:** None
- **Data normalization:** None

# Starting simple

## Steps

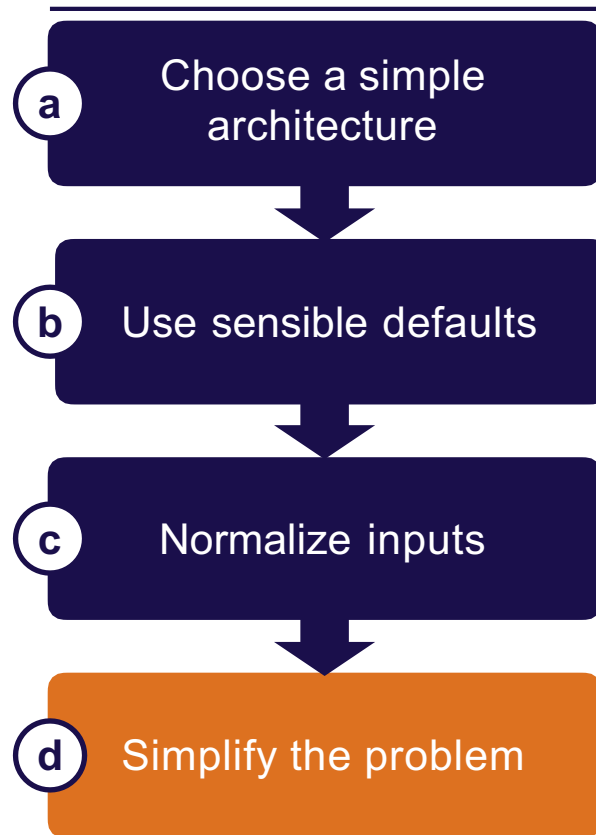


# Important to normalize scale of input data

- Subtract mean and divide by variance
- For images, fine to scale values to  $[0, 1]$  or  $[-0.5, 0.5]$  (e.g., by dividing by 255)  
[Careful, make sure your library doesn't do it for you!]

# Starting simple

## Steps



# Consider simplifying the problem as well

- Start with a small training set (~10,000 examples)
- Use a fixed number of objects, classes, image size, etc.
- Create a simpler synthetic training set



# Simplest model for pedestrian detection

- Start with a subset of 10,000 images for training, 1,000 for val, and 500 for test
- Use a LeNet architecture with sigmoid cross-entropy loss
- Adam optimizer with LR  $3e-4$
- No regularization

# Simplest model for pedestrian detection

- Start with a subset of 10,000 images for training and 10,000 for test
- Use a LeNet architecture with sigmoid cross-entropy loss
- Adam optimizer with LR  $3e-4$
- No regularization

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

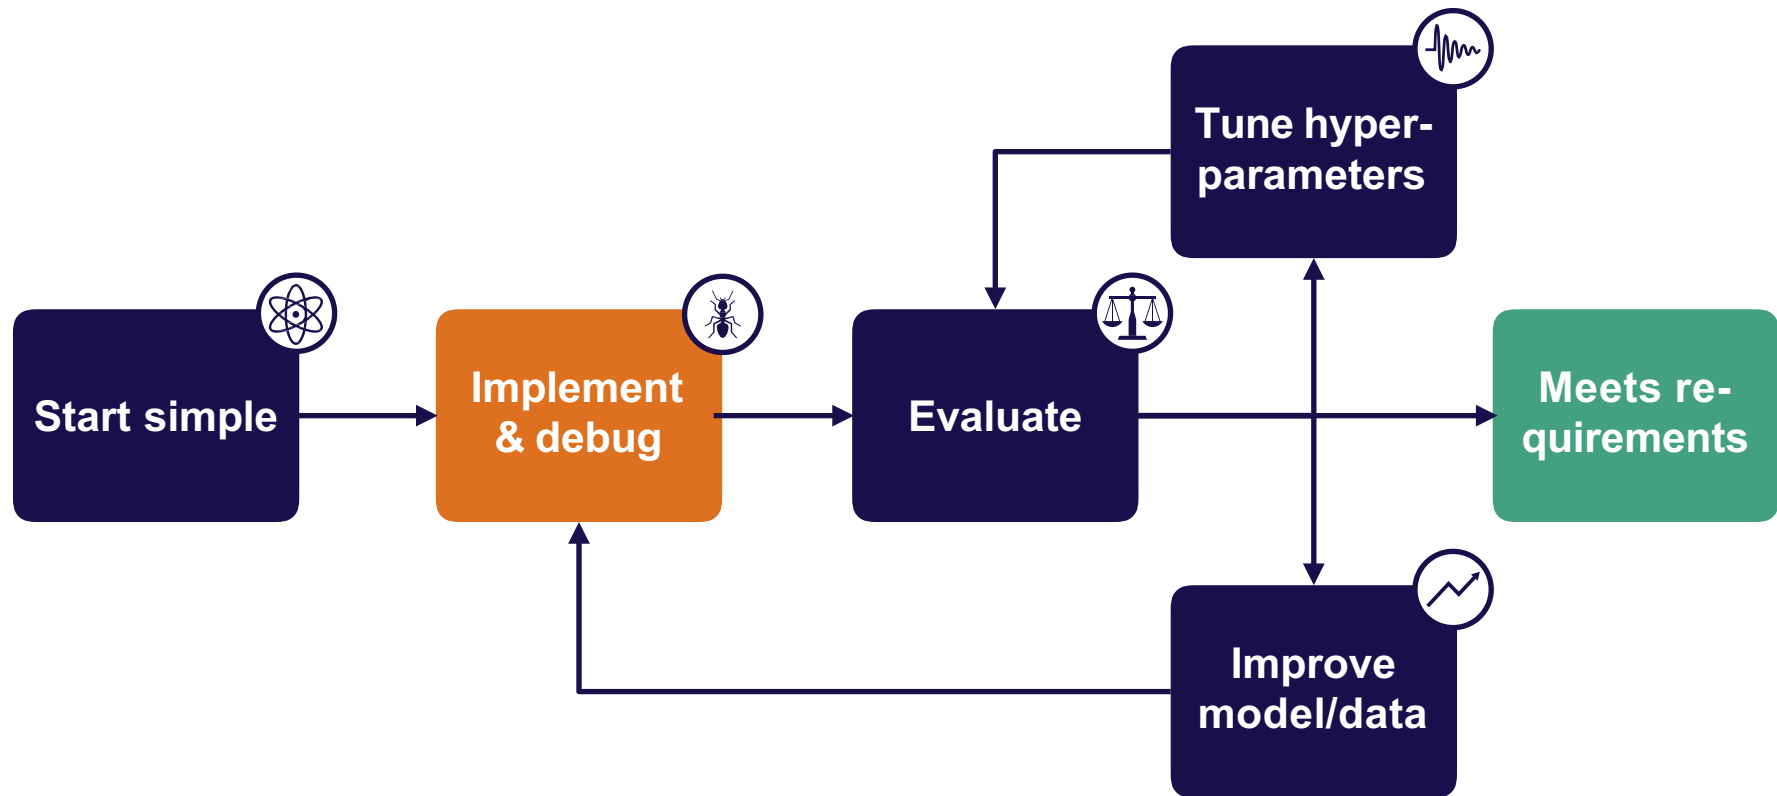
# Starting simple

Steps	Summary
<b>a</b> Choose a simple architecture	<ul style="list-style-type: none"><li>• LeNet, LSTM, or fully connected</li></ul>
<b>b</b> Use sensible defaults	<ul style="list-style-type: none"><li>• Adam optimizer &amp; no regularization</li></ul>
<b>c</b> Normalize inputs	<ul style="list-style-type: none"><li>• Subtract mean and divide by std, or just divide by 255 (ims)</li></ul>
<b>d</b> Simplify the problem	<ul style="list-style-type: none"><li>• Start with a simpler version of your problem (e.g., smaller dataset)</li></ul>

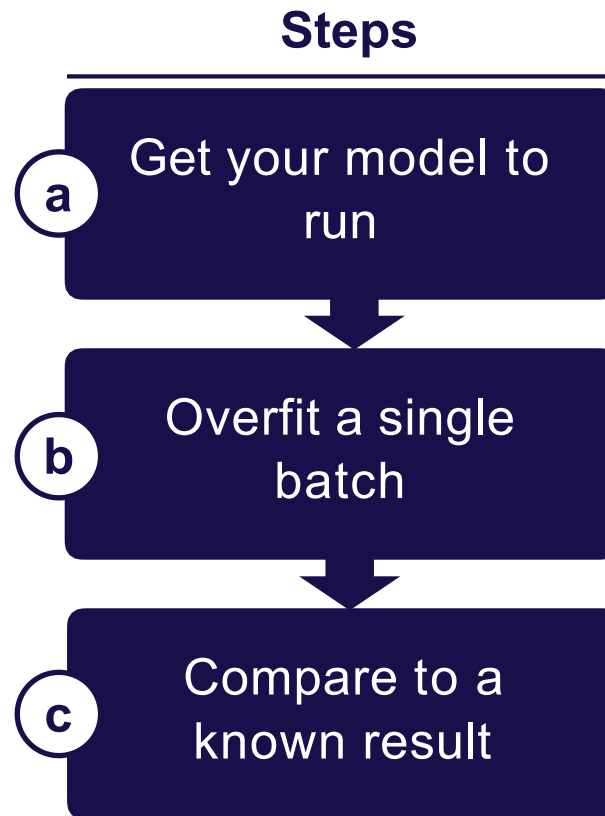


# Questions?

# Strategy for DL troubleshooting



# Implementing bug-free DL models



# Preview: the five most common DL bugs

- **Incorrect shapes for your tensors**

Can fail silently! E.g., accidental broadcasting: `x.shape = (None,)`, `y.shape = (None, 1)`, `(x+y).shape = (None, None)`



# Preview: the five most common DL bugs

- **Incorrect shapes for your tensors**

Can fail silently! E.g., accidental broadcasting: `x.shape = (None,)`, `y.shape = (None, 1)`, `(x+y).shape = (None, None)`

- **Pre-processing inputs incorrectly**

E.g., Forgetting to normalize, or too much pre-processing



# Preview: the five most common DL bugs

- **Incorrect shapes for your tensors**

Can fail silently! E.g., accidental broadcasting: `x.shape = (None,)`, `y.shape = (None, 1)`, `(x+y).shape = (None, None)`

- **Pre-processing inputs incorrectly**

E.g., Forgetting to normalize, or too much pre-processing

- **Incorrect input to your loss function**

E.g., softmaxed outputs to a loss that expects logits



# Preview: the five most common DL bugs

- **Incorrect shapes for your tensors**

Can fail silently! E.g., accidental broadcasting: `x.shape = (None,)`, `y.shape = (None, 1)`, `(x+y).shape = (None, None)`

- **Pre-processing inputs incorrectly**

E.g., Forgetting to normalize, or too much pre-processing

- **Incorrect input to your loss function**

E.g., softmaxed outputs to a loss that expects logits

- **Forgot to set up train mode for the net correctly**

E.g., toggling train/eval, controlling batch norm dependencies



# Preview: the five most common DL bugs

- **Incorrect shapes for your tensors**

Can fail silently! E.g., accidental broadcasting: `x.shape = (None,)`, `y.shape = (None, 1)`, `(x+y).shape = (None, None)`

- **Pre-processing inputs incorrectly**

E.g., Forgetting to normalize, or too much pre-processing

- **Incorrect input to your loss function**

E.g., softmaxed outputs to a loss that expects logits

- **Forgot to set up train mode for the net correctly**

E.g., toggling train/eval, controlling batch norm dependencies

- **Numerical instability - inf/NaN**

Often stems from using an exp, log, or div operation





# General advice for implementing your model

## Lightweight implementation

- Minimum possible new lines of code for v1
- Rule of thumb: <200 lines
- (Tested infrastructure components are fine)

## Use off-the-shelf components, e.g.,

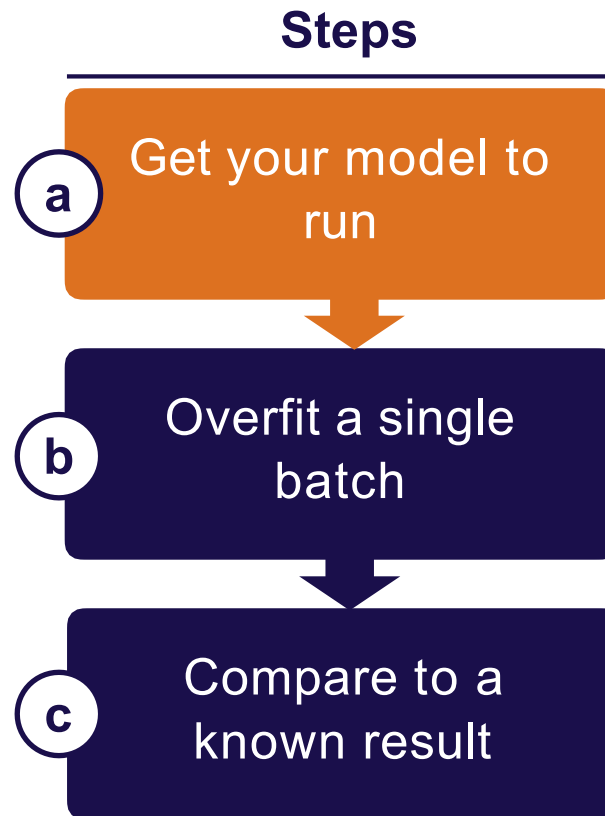
- Keras
- `tf.layers.dense(...)`  
instead of  
`tf.nn.relu(tf.matmul(W, x))`
- `tf.losses.cross_entropy(...)`  
instead of writing out the exp

## Build complicated data pipelines later

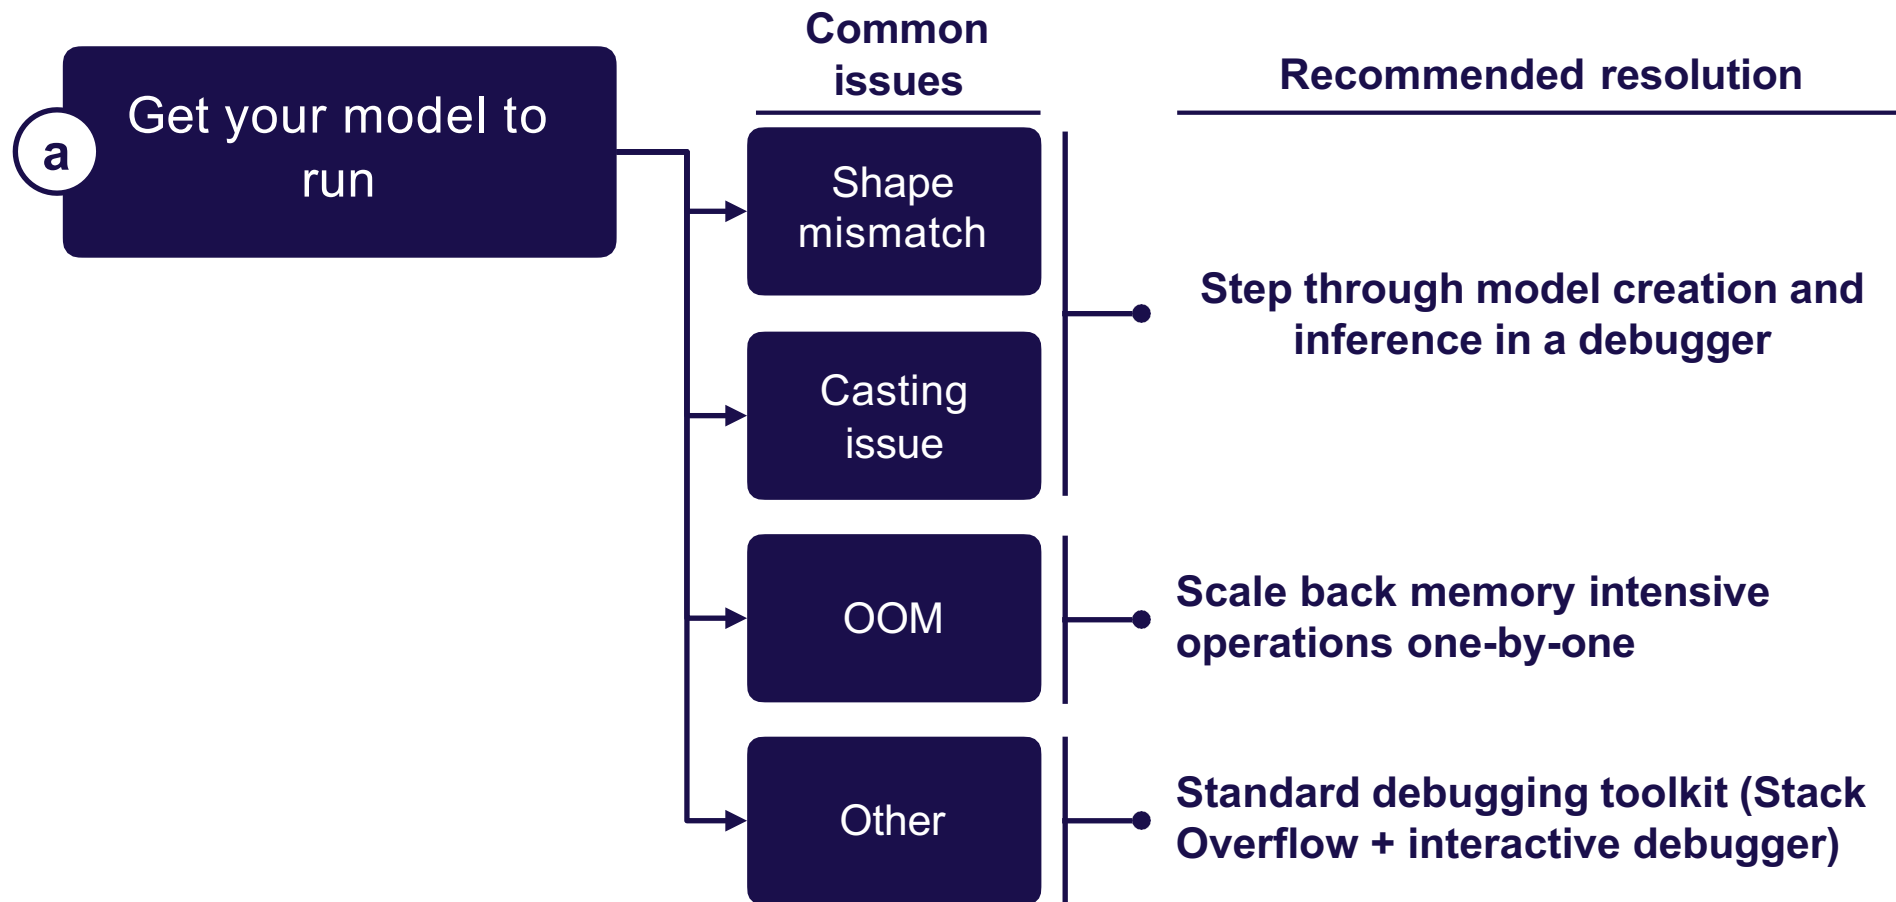
- Start with a dataset you can load into memory



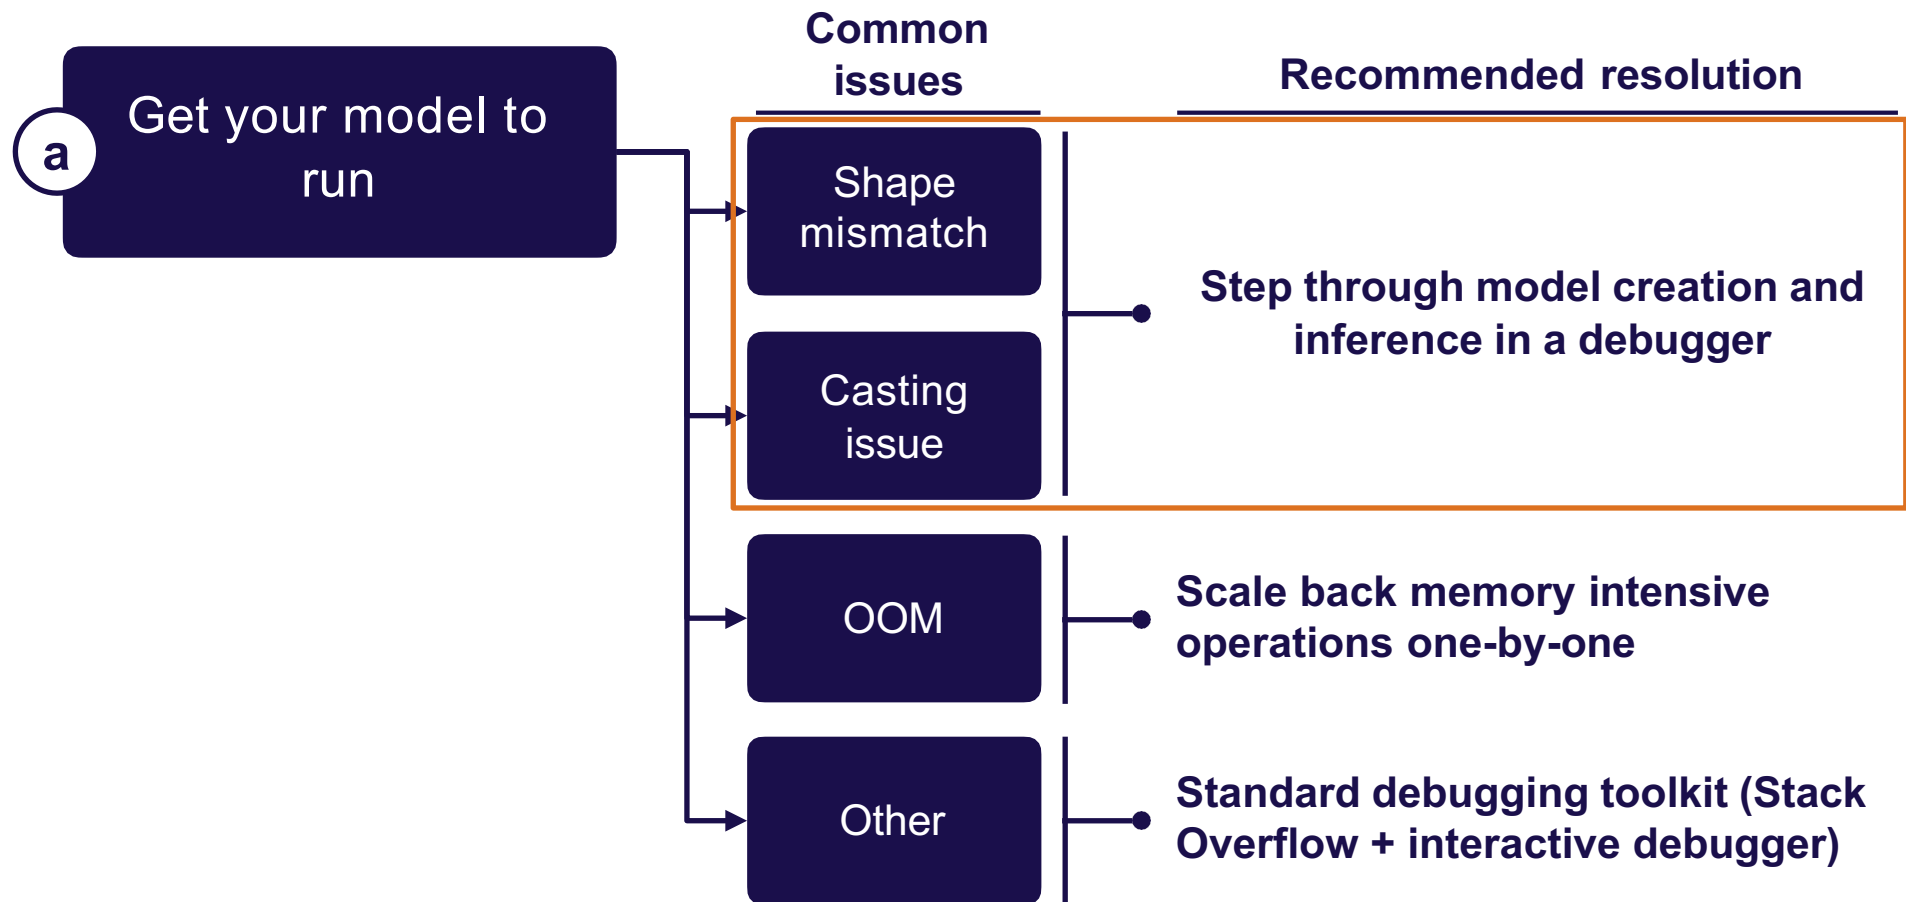
# Implementing bug-free DL models



# Implementing bug-free DL models



# Implementing bug-free DL models



# Debuggers for DL code

- Pytorch: easy, use ipdb
- tensorflow: trickier

## Option 1: step through graph creation

```
2 # Option 1: step through graph creation
3 import ipdb; ipdb.set_trace()
4
5 for i in range(num_layers):
6     out = layers.fully_connected(out, 50)
7
```

```
josh at MacBook-Pro-9 in ~/projects
$ python test.py
> /Users/josh/projects/test.py(5)<module>()
   3 h = tf.placeholder(tf.float32, (None, 100))
   4 import ipdb; ipdb.set_trace()
----> 5 w = tf.layers.dense(h)

ipdb> 
```



# Debuggers for DL code

- Pytorch: easy, use ipdb
- tensorflow: trickier

## Option 2: step into training loop

```
9 # Option 2: step into training loop
10 sess = tf.Session()
11 for i in range(num_epochs):
12     import ipdb; ipdb.set_trace()
13     loss_, _ = sess.run([loss, train_op])
14
```



**Evaluate tensors using sess.run(...)**



# Debuggers for DL code

- Pytorch: easy, use ipdb
- tensorflow: trickier

## Option 3: use tfdbg

```
python -m tensorflow.python.debug.examples.debug_mnist --debug
```


```
-- run-start: run #1: 1 fetch (accuracy/accuracy/Mean:0); 2 feeds -----
| <- -> | run_info
| run | invoke_stepper | exit |
TTTTTT FFFF DDD BBBB GGG
TT F D D B B G
TT FFF D D BBBB G GG
TT F D D B B G G
TT F DDD BBBB GGG

=====
Session.run() call #1:
Fetch(es):
  accuracy/accuracy/Mean:0
Feed dict(s):
  input/x-input:0
  input/y-input:0
=====
Select one of the following commands to proceed ---->
run:
  Execute the run() call with debug tensor-watching
run -n:
  Execute the run() call without debug tensor-watching
-- Scroll (PgDn): 0.00% ----- Mouse: ON --
tfdbg> |
```

**Stops  
execution at  
each  
sess.run(...)  
and lets you  
inspect**

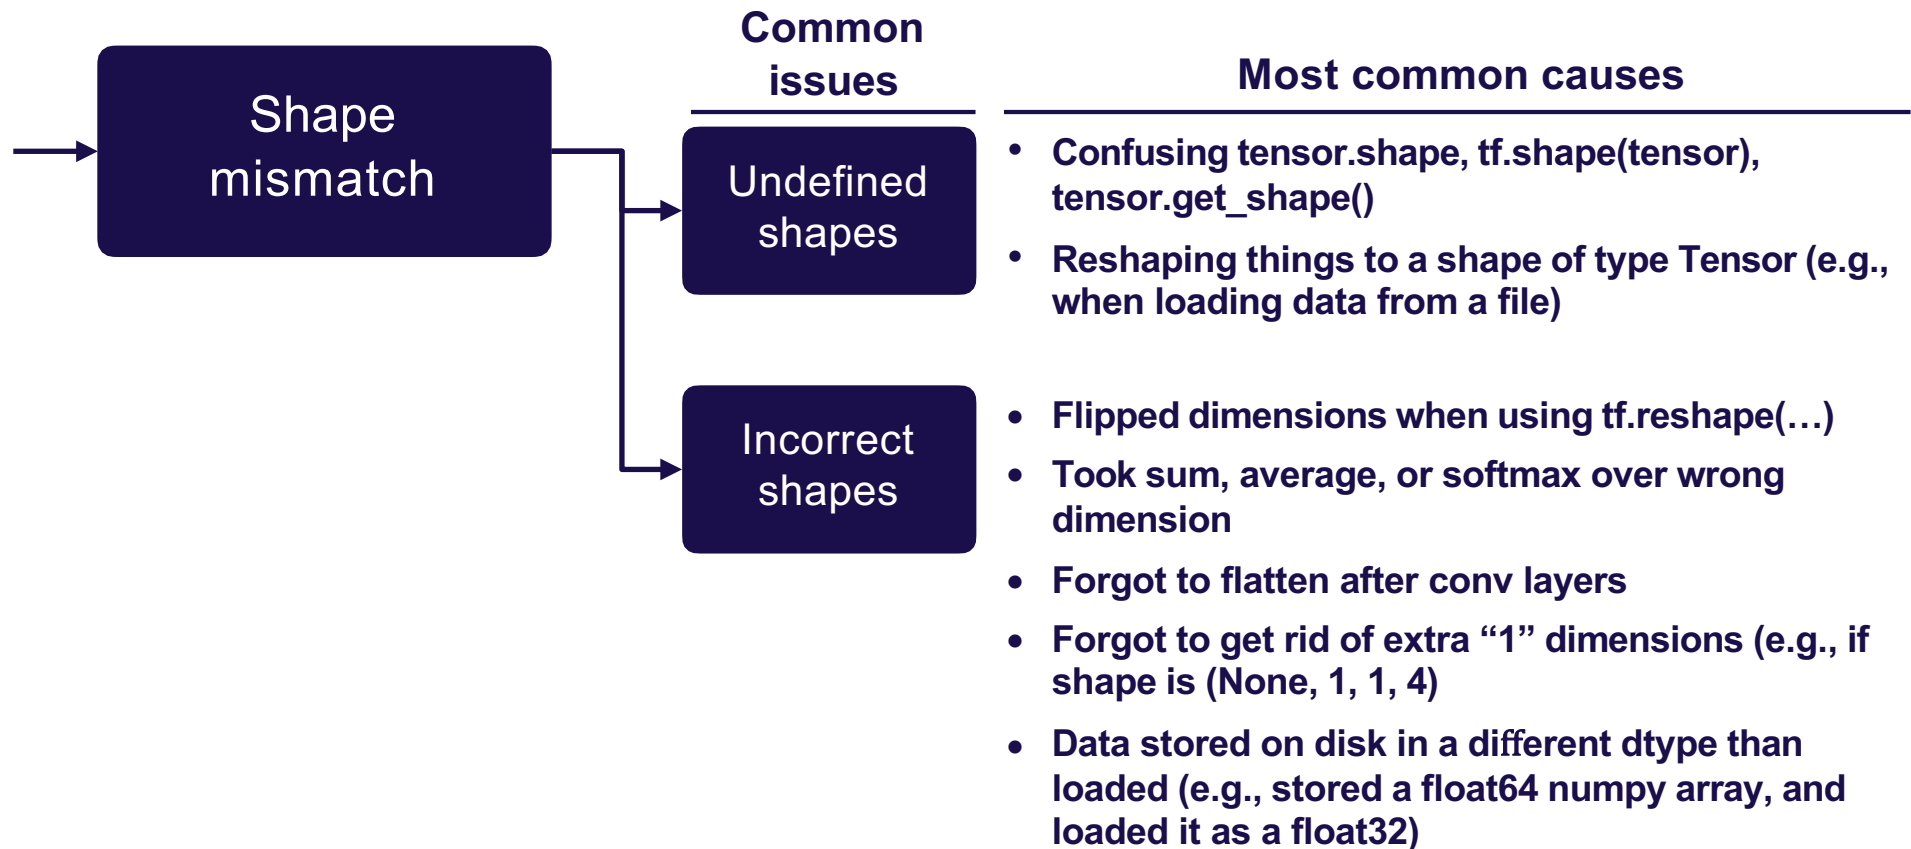


# Implementing bug-free DL models






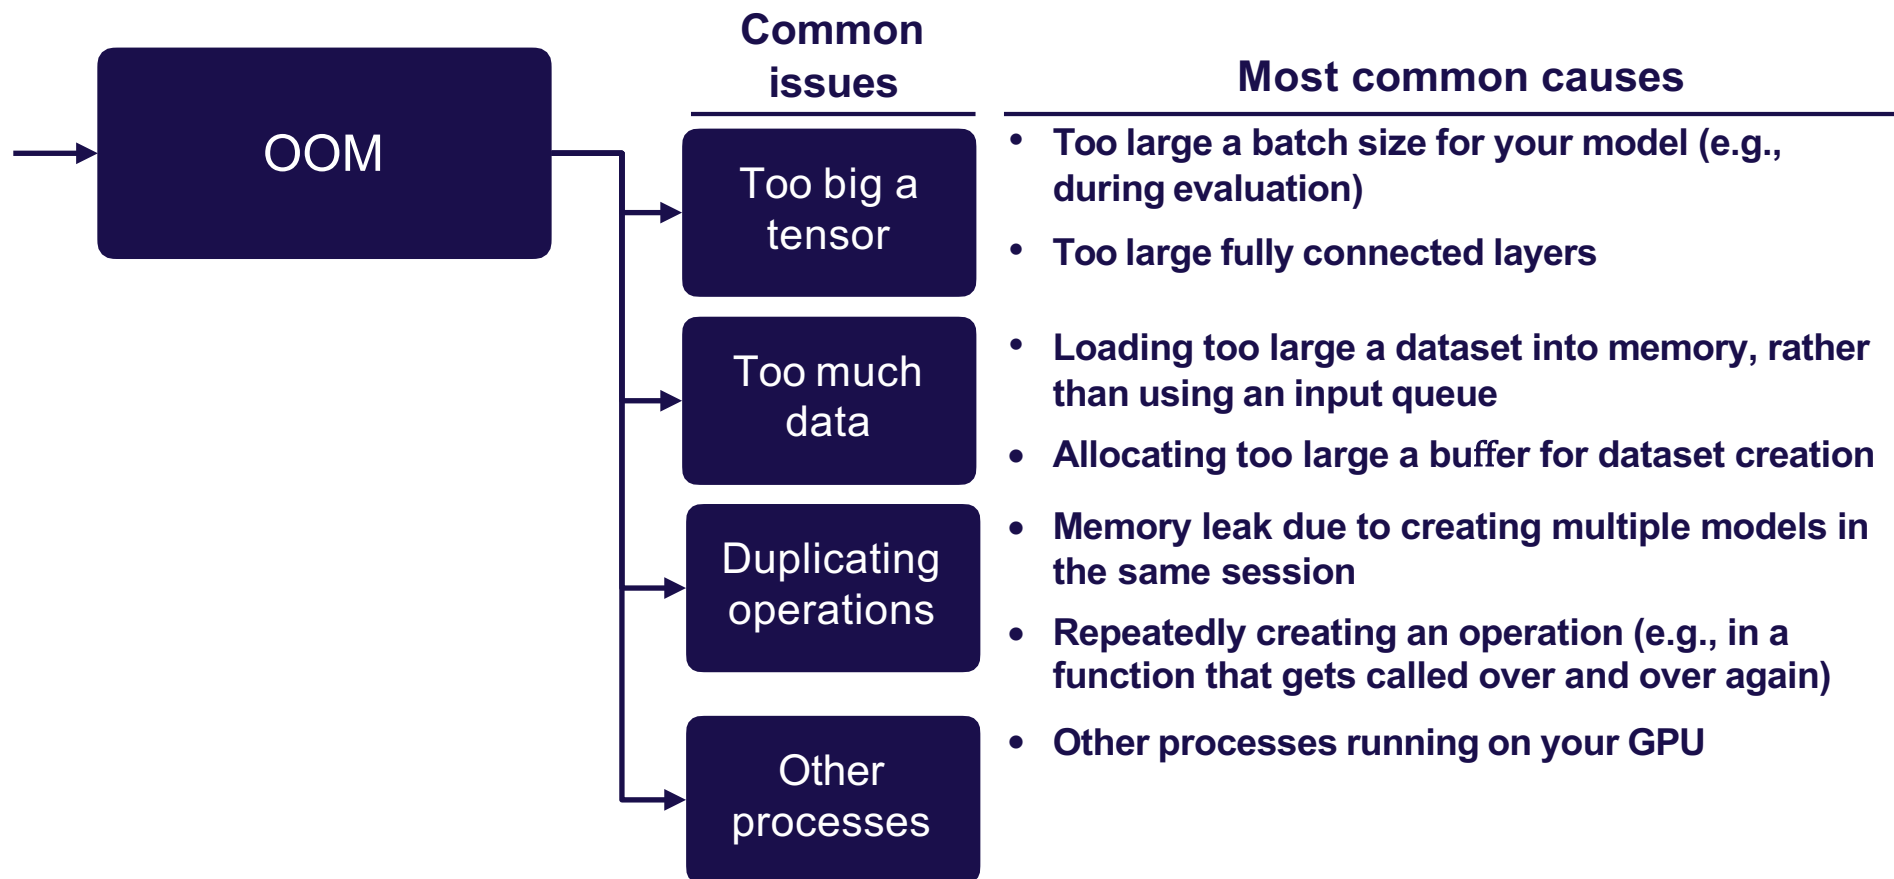
# Implementing bug-free DL models



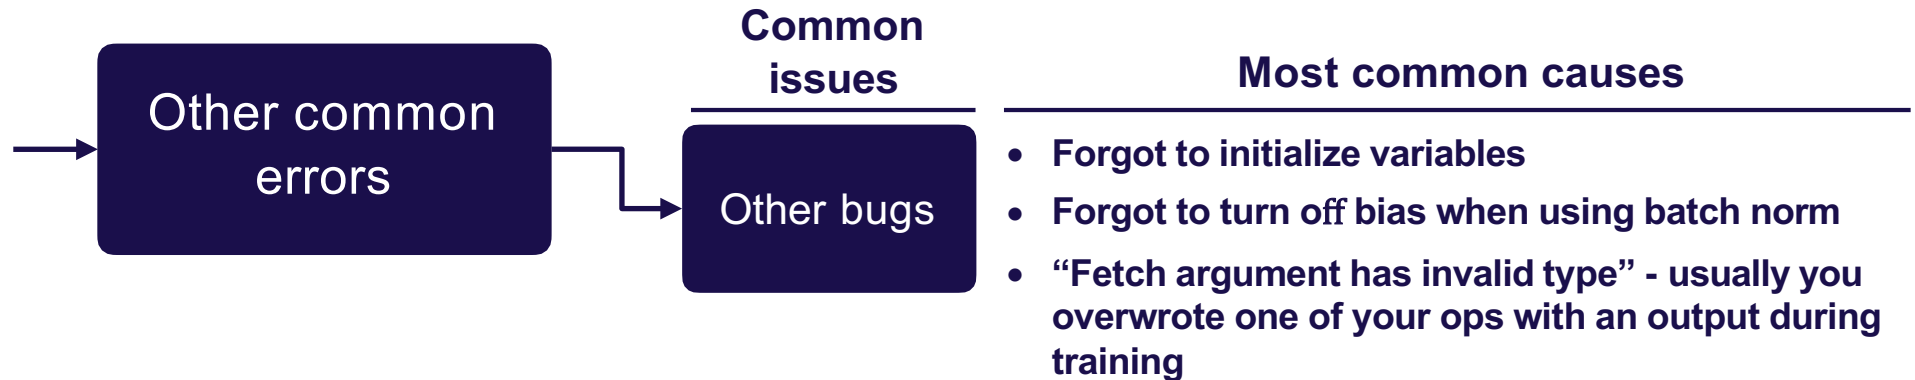
# Implementing bug-free DL models



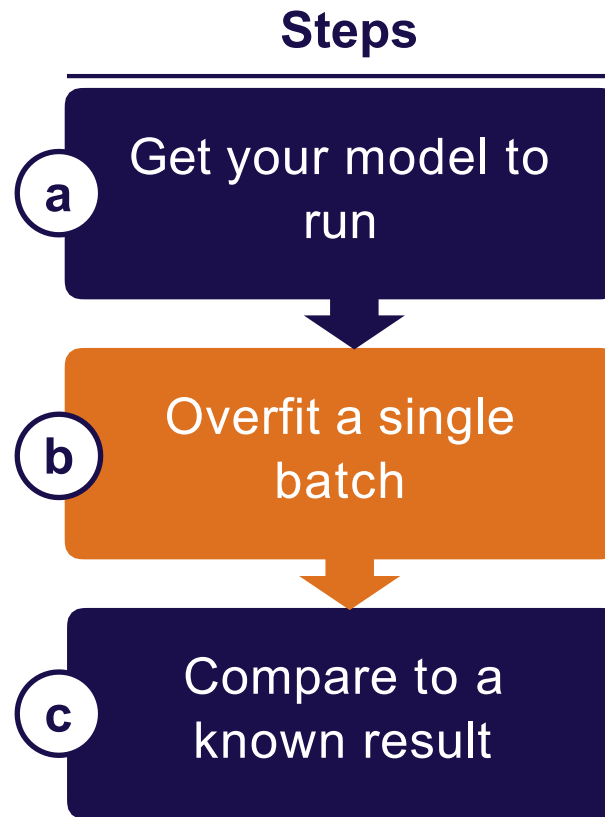
# Implementing bug-free DL models



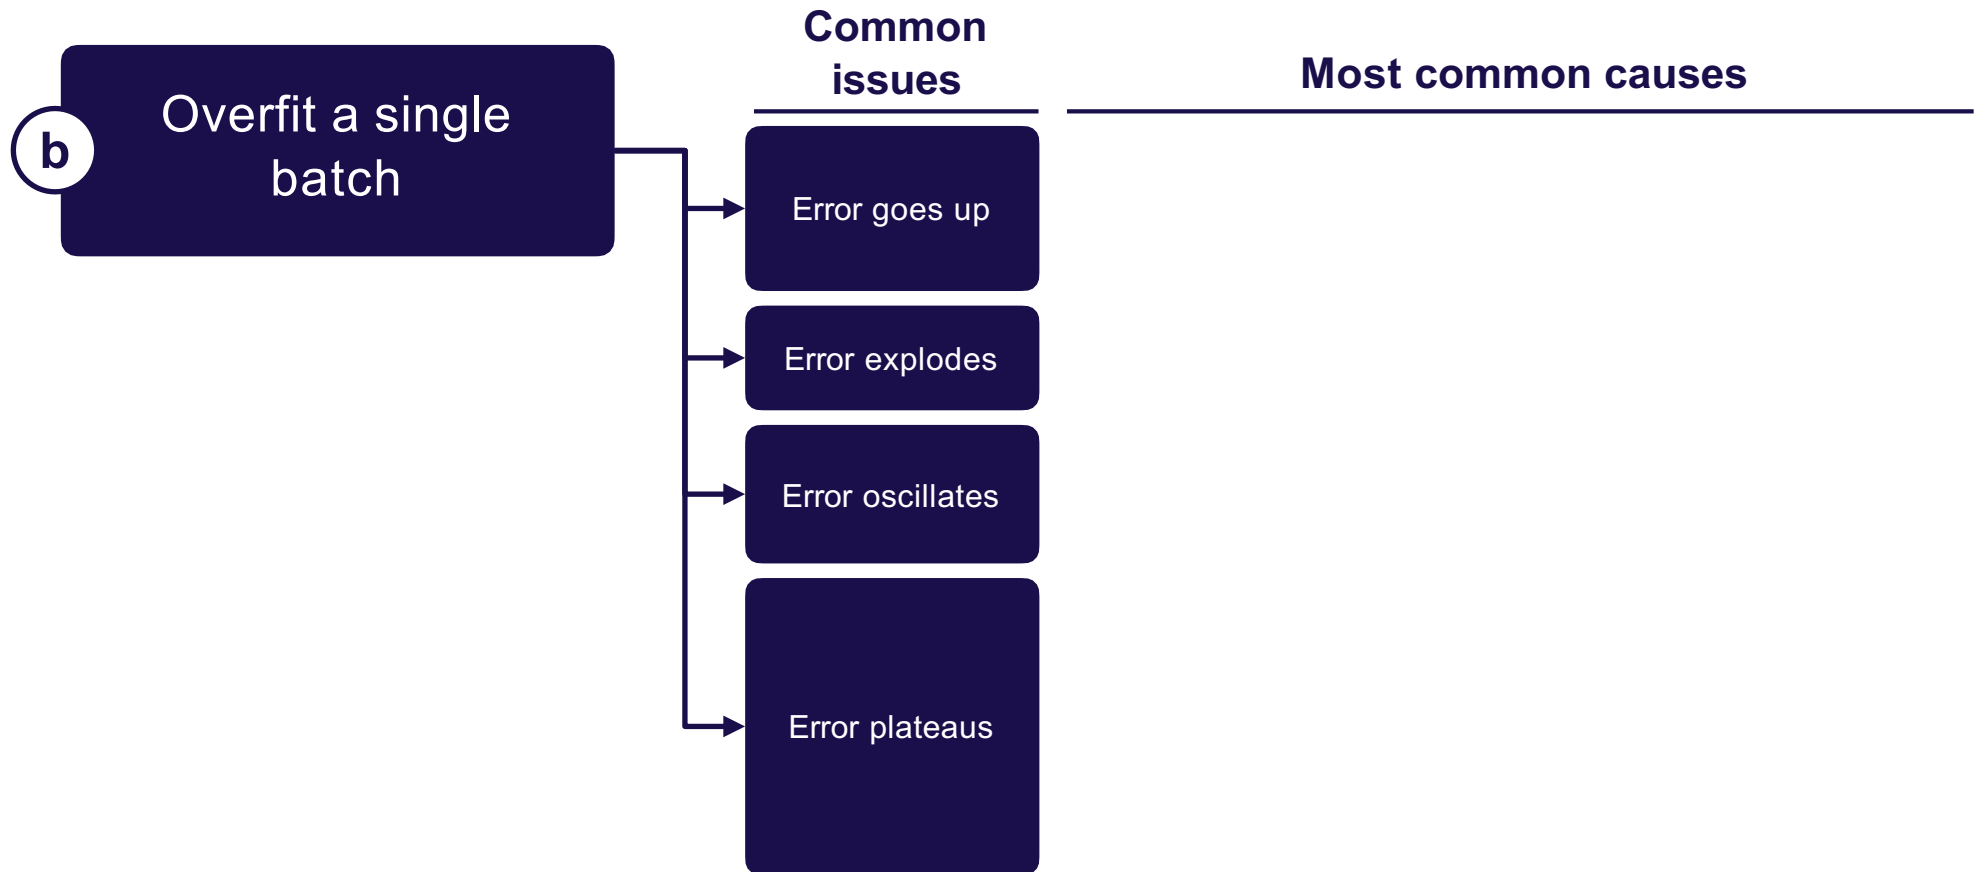
# Implementing bug-free DL models




# Implementing bug-free DL models



# Implementing bug-free DL models




# Implementing bug-free DL models



<https://medium.com/data-from-the-trenches/the-learning-rate-black-magic-c4a652133cd7>



# Implementing bug-free DL models




<https://medium.com/data-from-the-trenches/the-learning-rate-black-magic-c4a652133cd7>





# Implementing bug-free DL models



**Error oscillates**

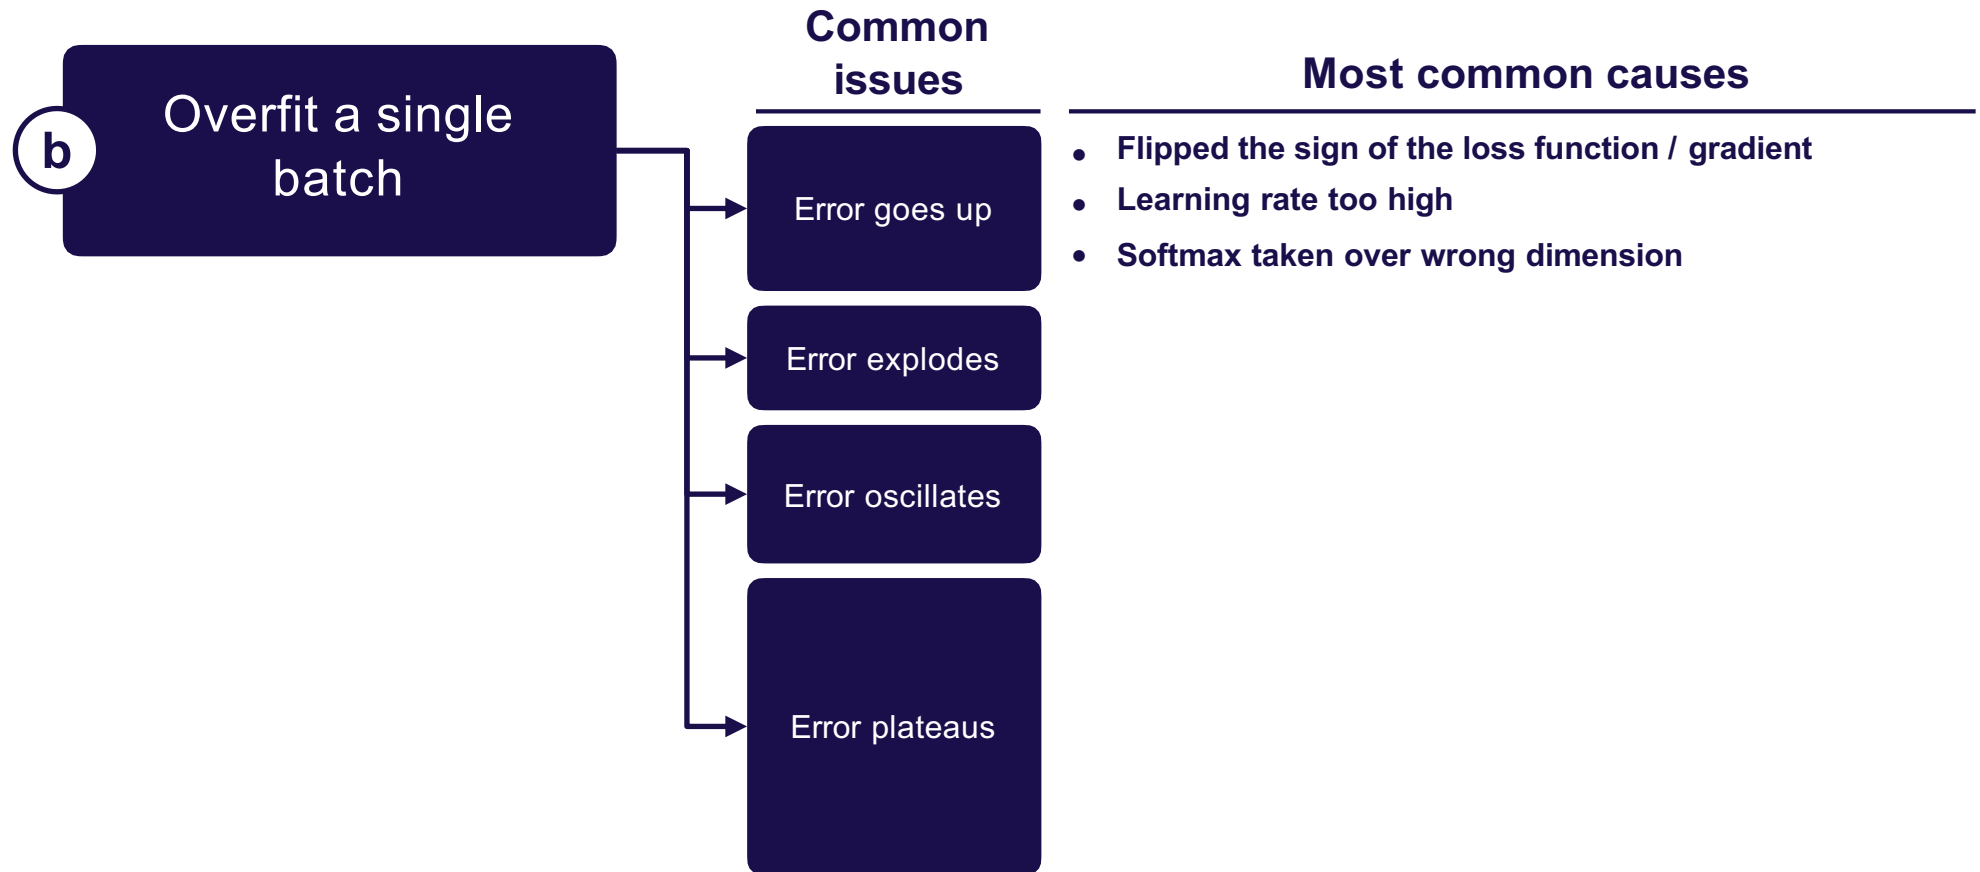
<https://medium.com/data-from-the-trenches/the-learning-rate-black-magic-c4a652133cd7>



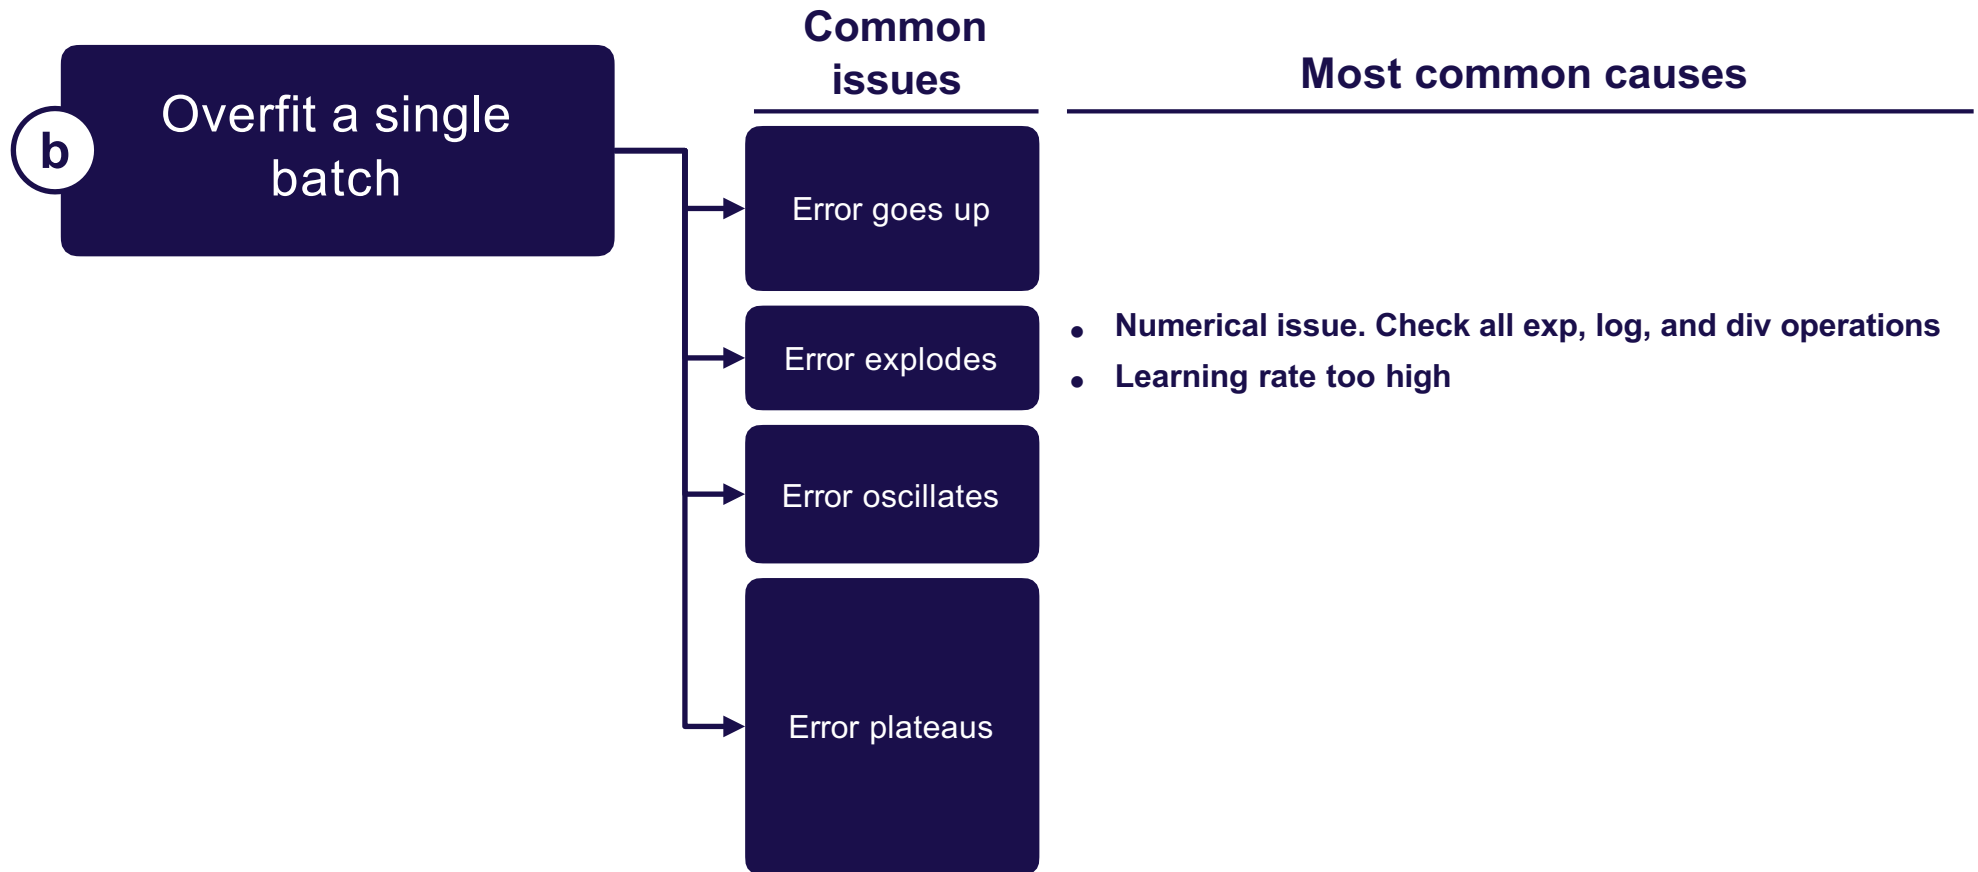
# Implementing bug-free DL models



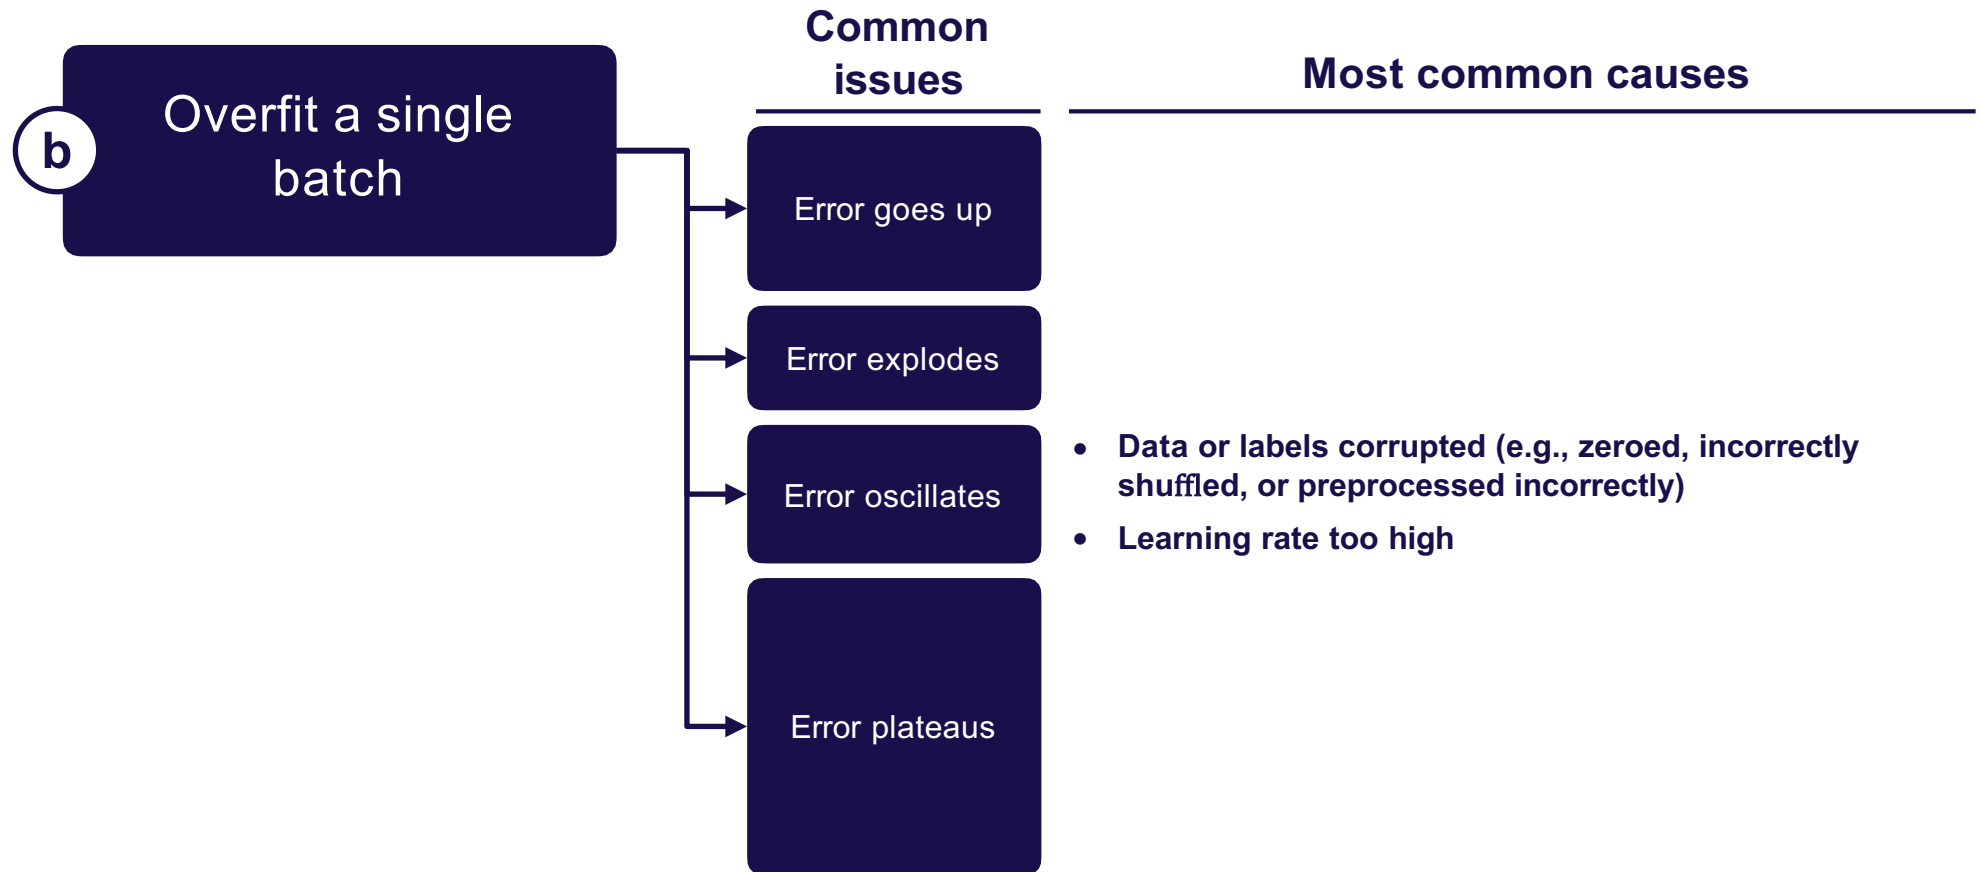
# Implementing bug-free DL models



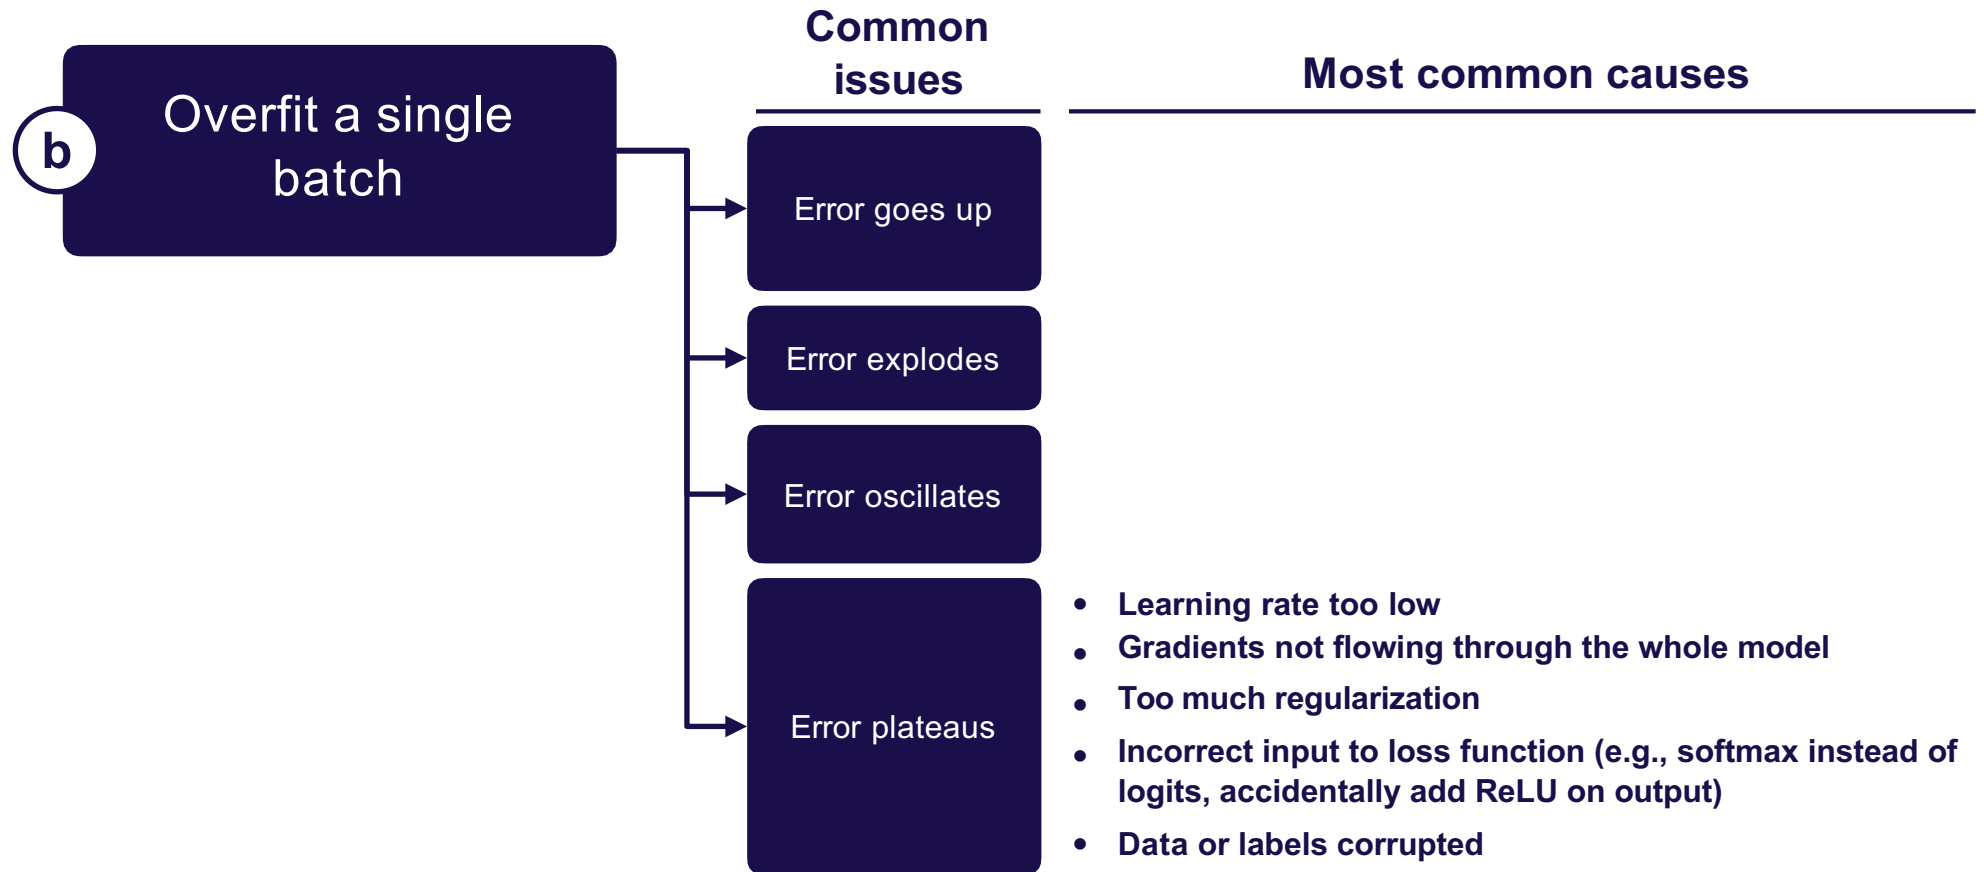
# Implementing bug-free DL models



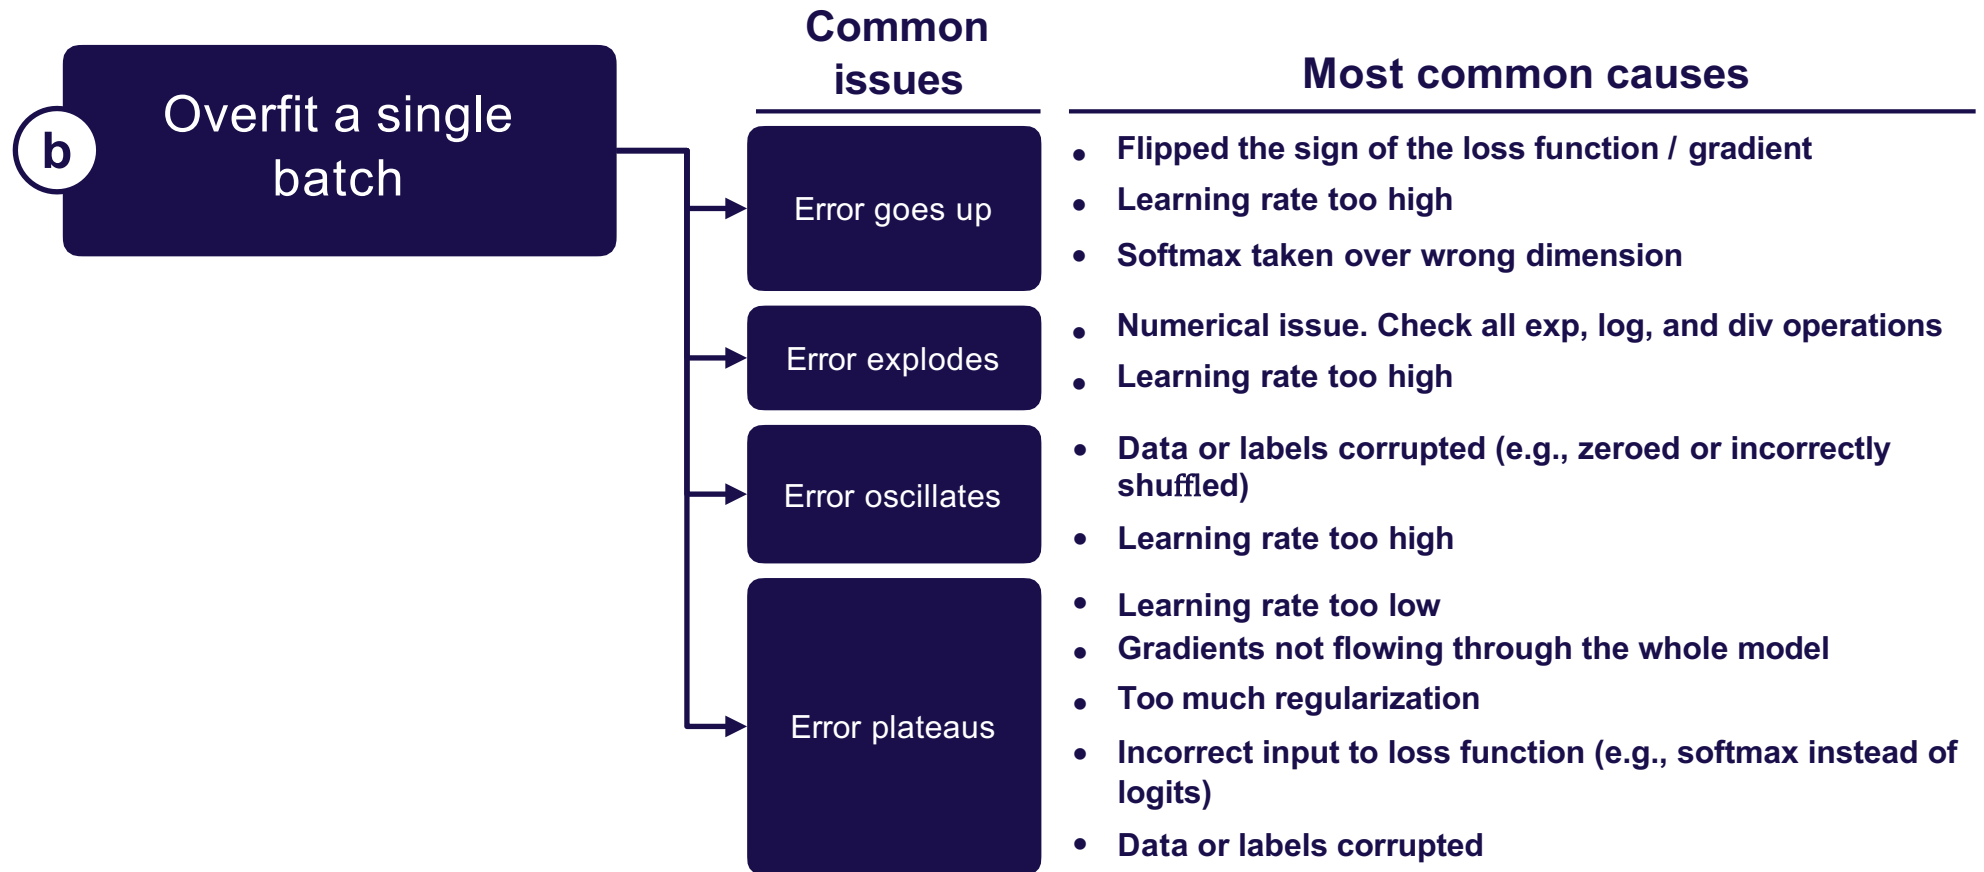
# Implementing bug-free DL models



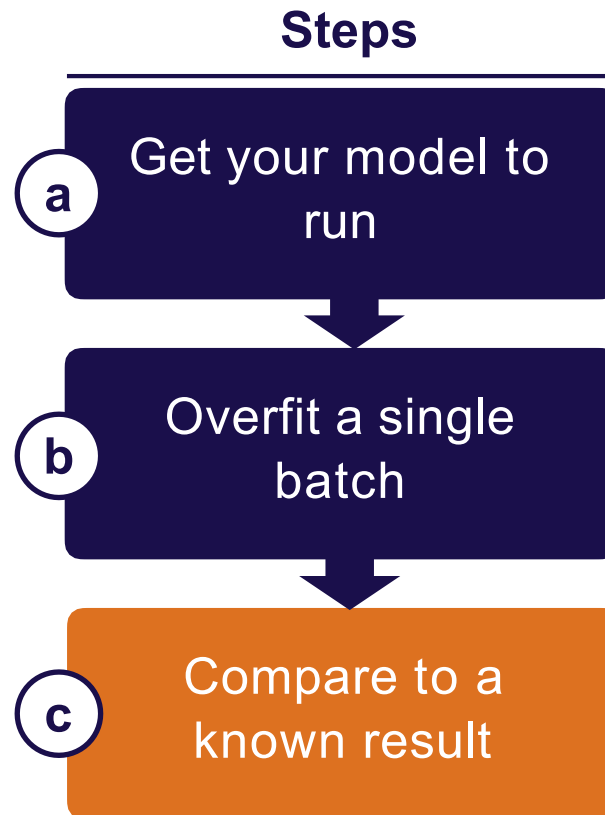
# Implementing bug-free DL models



# Implementing bug-free DL models



# Implementing bug-free DL models





# Hierarchy of known results

**More  
useful**

- Official model implementation evaluated on similar dataset to yours

## **You can:**

- Walk through code line-by-line and ensure you have the same output
- Ensure your performance is up to par with expectations

**Less  
useful**



# Hierarchy of known results

**More  
useful**

- Official model implementation evaluated on benchmark (e.g., MNIST)

**You can:**

- Walk through code line-by-line and ensure you have the same output

**Less  
useful**



# Hierarchy of known results

**More  
useful**

- Unofficial model implementation

**You can:**

- Same as before, but with lower confidence

**Less  
useful**



# Hierarchy of known results

**More  
useful**

- Results from a paper (with no code)

**You can:**

- Ensure your performance is up to par with expectations

**Less  
useful**



# Hierarchy of known results

**More  
useful**

**You can:**

- Make sure your model performs well in a simpler setting

- Results from your model on a benchmark dataset (e.g., MNIST)

**Less  
useful**



# Hierarchy of known results

**More  
useful**

**Less  
useful**

**You can:**

- Get a general sense of what kind of performance can be expected
- 
- Results from a similar model on a similar dataset



# Hierarchy of known results

**More  
useful**

**You can:**

- Make sure your model is learning anything at all

**Less  
useful**

- Super simple baselines (e.g., average of outputs or linear regression)



# Hierarchy of known results

**More  
useful**

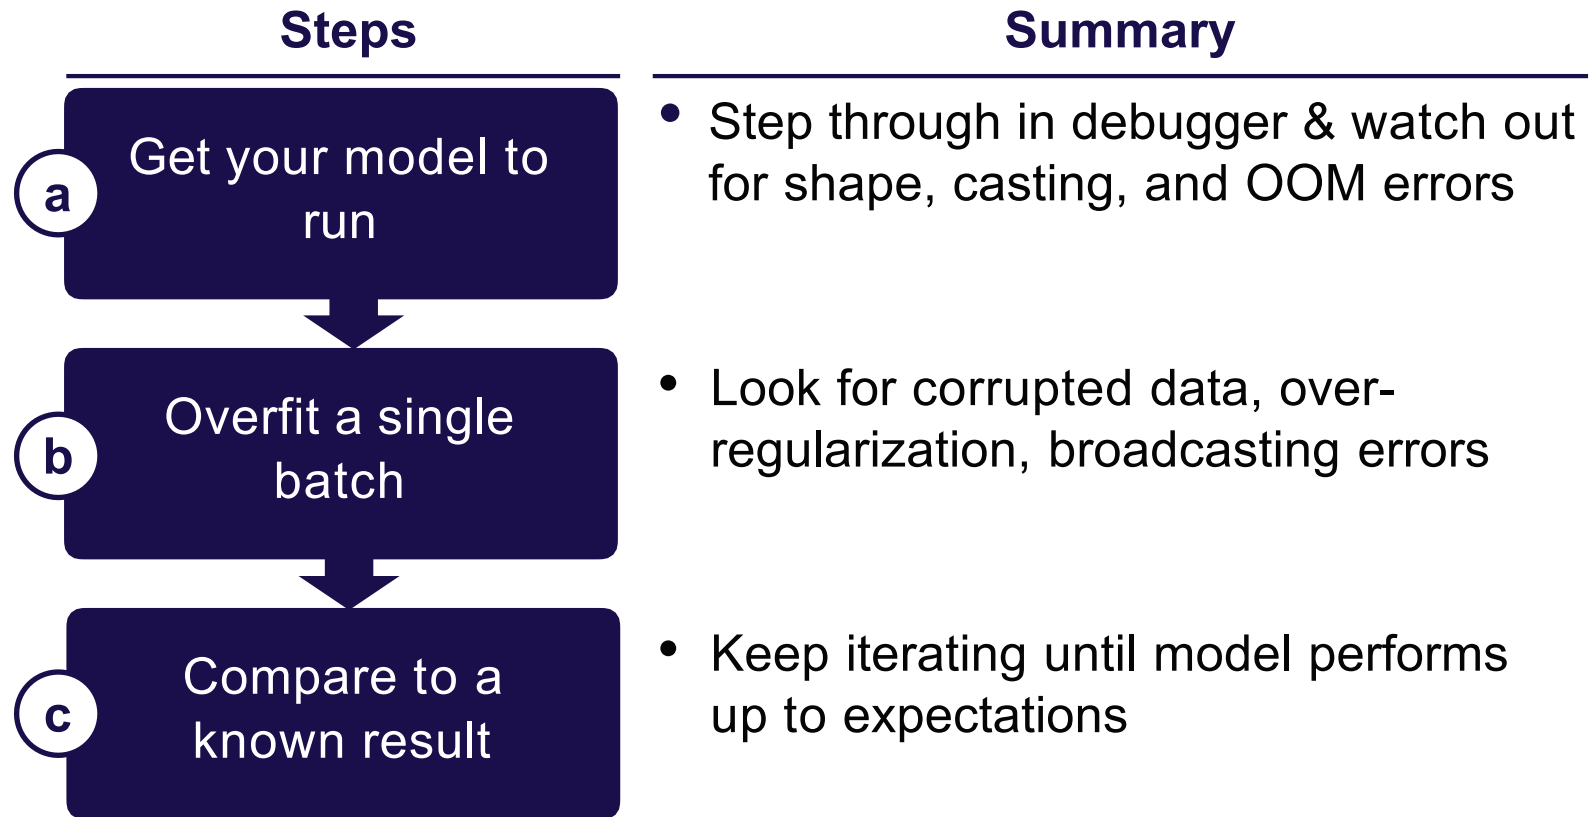
- Official model implementation evaluated on similar dataset to yours
- Official model implementation evaluated on benchmark (e.g., MNIST)
- Unofficial model implementation
- Results from the paper (with no code)
- Results from your model on a benchmark dataset (e.g., MNIST)
- Results from a similar model on a similar dataset
- Super simple baselines (e.g., average of outputs or linear regression)

**Less  
useful**





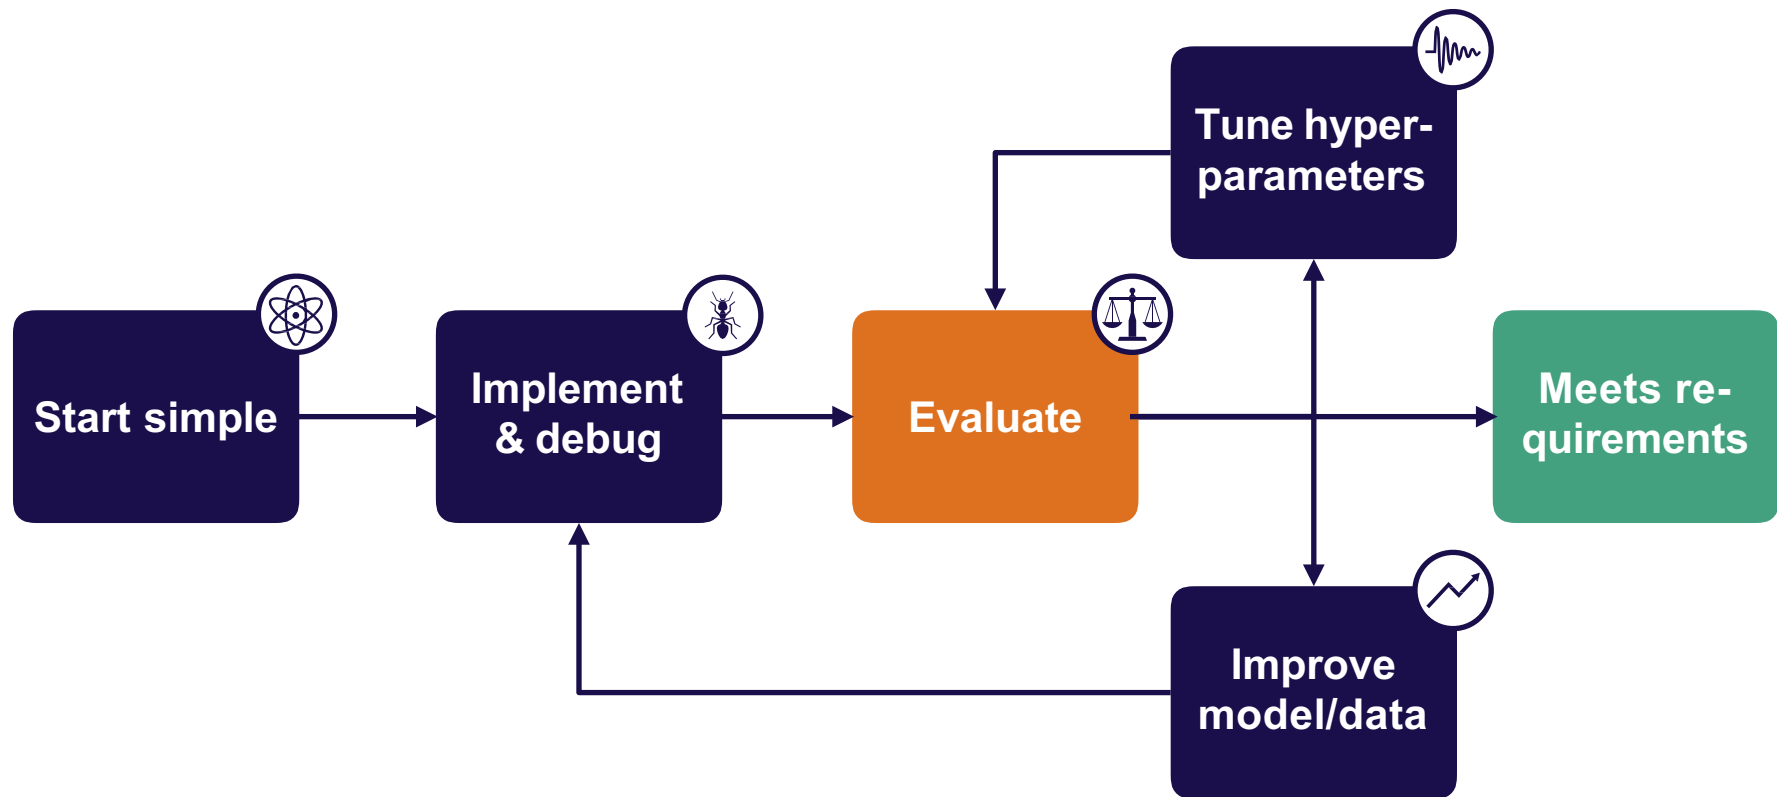
# Summary: how to implement & debug




# Questions?




# Strategy for DL troubleshooting




# Bias-variance decomposition





# Bias-variance decomposition



# Bias-variance decomposition



# Bias-variance decomposition



# Bias-variance decomposition

- **Test error = irreducible error + bias + variance + val overfitting**
- This assumes train, val, and test all come from the same distribution. What if not?





# Handling distribution shift

**Train data**




**Test data**




**Use two val sets: one sampled from training distribution and one from test distribution**



# The bias-variance tradeoff



# Bias-variance with distribution shift



# Bias-variance with distribution shift



# Train, val, and test error for pedestrian detection

Error source	Value
Goal performance	1%
Train error	20%
Validation error	27%
Test error	28%

Train - goal = 19%  
(under-fitting)

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Error source	Value
Goal performance	1%
Train error	20%
Validation error	27%
Test error	28%

Val - train = 7%  
(over-fitting)

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Error source	Value
Goal performance	1%
Train error	20%
Validation error	27%
Test error	28%

Test - val = 1%  
(looks good!)

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy



# Summary: evaluating model performance

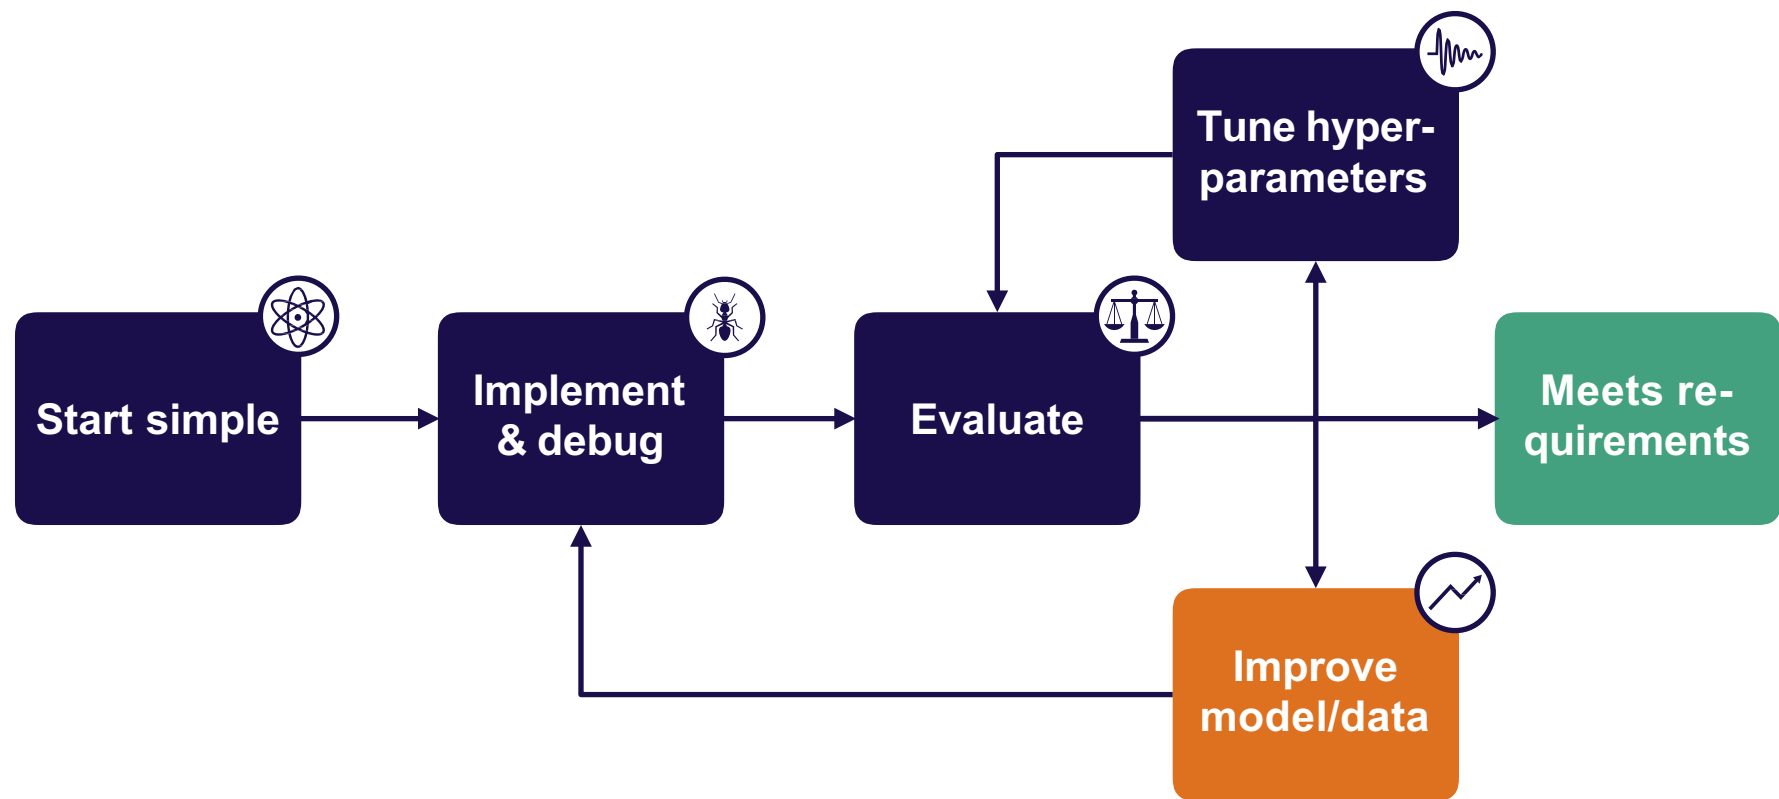
**Test error = irreducible error + bias + variance  
+ distribution shift + val overfitting**



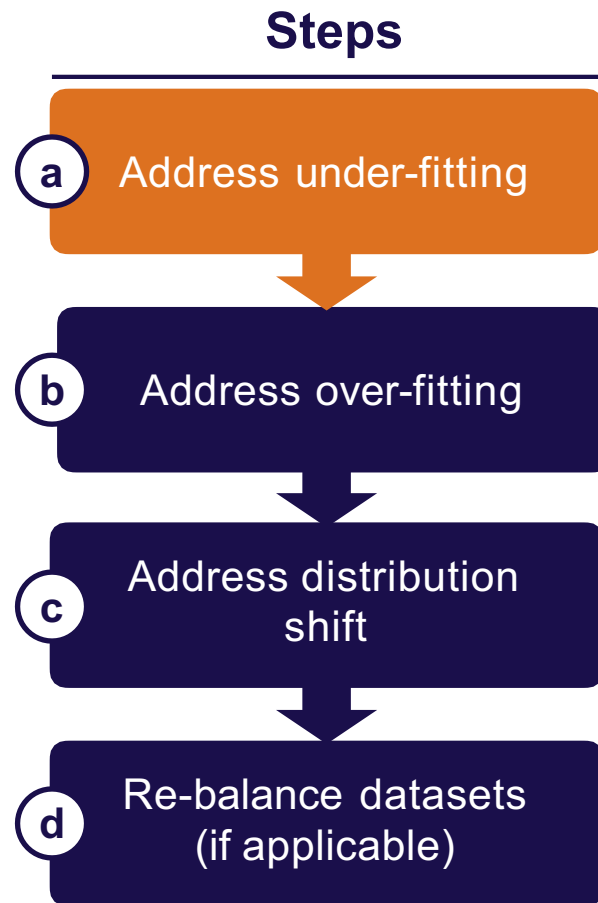


# Questions?

# Strategy for DL troubleshooting




# Prioritizing improvements (i.e., applied b-v)



# Addressing under-fitting (i.e., reducing bias)

**Try first**

- 
- A. Make your model bigger (i.e., add layers or use more units per layer)
  - B. Reduce regularization
  - C. Error analysis
  - D. Choose a different (closer to state-of-the art) model architecture (e.g., move from LeNet to ResNet)
  - E. Tune hyper-parameters (e.g., learning rate)

**Try later**

- F. Add features

# Train, val, and test error for pedestrian detection

Add more layers  
to the ConvNet



Error source	Value	Value
Goal performance	1%	1%
Train error	20%	7%
Validation error	27%	19%
Test error	28%	20%



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy  
(i.e., 1% error)

# Train, val, and test error for pedestrian detection

Switch to  
ResNet-101



Error source	Value	Value	Value
Goal performance	<del>1%</del>	<del>1%</del>	1%
Train error	<del>20%</del>	<del>7%</del>	3%
Validation error	<del>27%</del>	<del>19%</del>	10%
Test error	<del>28%</del>	<del>20%</del>	10%



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy  
(i.e., 1% error)

# Train, val, and test error for pedestrian detection

Tune learning  
rate



Error source	Value	Value	Value	Value
Goal performance	1%	1%	1%	1%
Train error	20%	7%	3%	0.8%
Validation error	27%	19%	10%	12%
Test error	28%	20%	10%	12%

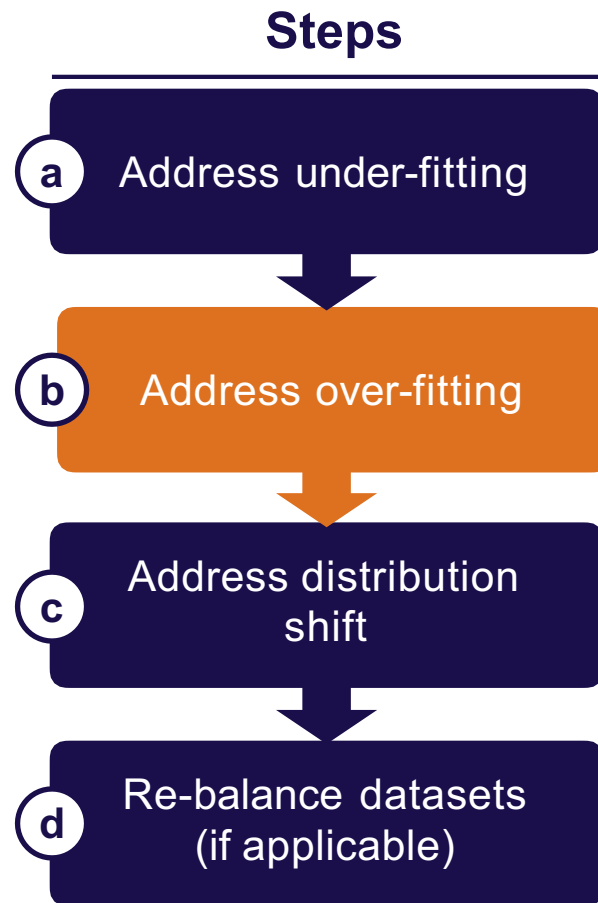


0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy  
(i.e., 1% error)


# Prioritizing improvements (i.e., applied b-v)





# Addressing over-fitting (i.e., reducing variance)

**Try first**

- 
- A. Add more training data (if possible!)
  - B. Add normalization (e.g., batch norm, layer norm)
  - C. Add data augmentation
  - D. Increase regularization (e.g., dropout, L2, weight decay)
  - E. Error analysis
  - F. Choose a different (closer to state-of-the-art) model architecture
  - G. Tune hyperparameters
  - H. Early stopping
  - I. Remove features
  - J. Reduce model size

**Try later**

# Addressing over-fitting (i.e., reducing variance)

**Try first**

- A. Add more training data (if possible!)
- B. Add normalization (e.g., batch norm, layer norm)
- C. Add data augmentation
- D. Increase regularization (e.g., dropout, L2, weight decay)
- E. Error analysis
- F. Choose a different (closer to state-of-the-art) model architecture
- G. Tune hyperparameters

**Try later**

- H. Early stopping
- I. Remove features
- J. Reduce model size

**Not  
recommended!**

# Train, val, and test error for pedestrian detection

Error source	Value
Goal performance	1%
Train error	0.8%
Validation error	12%
Test error	12%

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Increase dataset  
size to 250,000



Error source	Value	Value
Goal performance	<del>1%</del>	1%
Train error	<del>0.8%</del>	1.5%
Validation error	<del>12%</del>	5%
Test error	<del>12%</del>	6%

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Add weight  
decay



Error source	Value	Value	Value
Goal performance	1%	1%	1%
Train error	<del>0.8%</del>	<del>1.5%</del>	1.7%
Validation error	<del>12%</del>	5%	4%
Test error	<del>12%</del>	<del>6%</del>	4%

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Add data  
augmentation



Error source	Value	Value	Value	Value
Goal performance	1%	1%	1%	1%
Train error	<del>0.8%</del>	<del>1.5%</del>	<del>1.7%</del>	2%
Validation error	<del>12%</del>	<del>5%</del>	<del>4%</del>	2.5%
Test error	<del>12%</del>	<del>6%</del>	<del>4%</del>	2.6%

Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Train, val, and test error for pedestrian detection

Tune num layers, optimizer params, weight initialization, kernel size, weight decay



Error source	Value	Value	Value	Value	Value
Goal performance	1%	1%	1%	1%	1%
Train error	<del>0.8%</del>	<del>1.5%</del>	<del>1.7%</del>	<del>2%</del>	0.6%
Validation error	<del>12%</del>	<del>5%</del>	<del>4%</del>	<del>2.5%</del>	0.9%
Test error	<del>12%</del>	<del>6%</del>	<del>4%</del>	<del>2.6%</del>	1.0%

## Running example

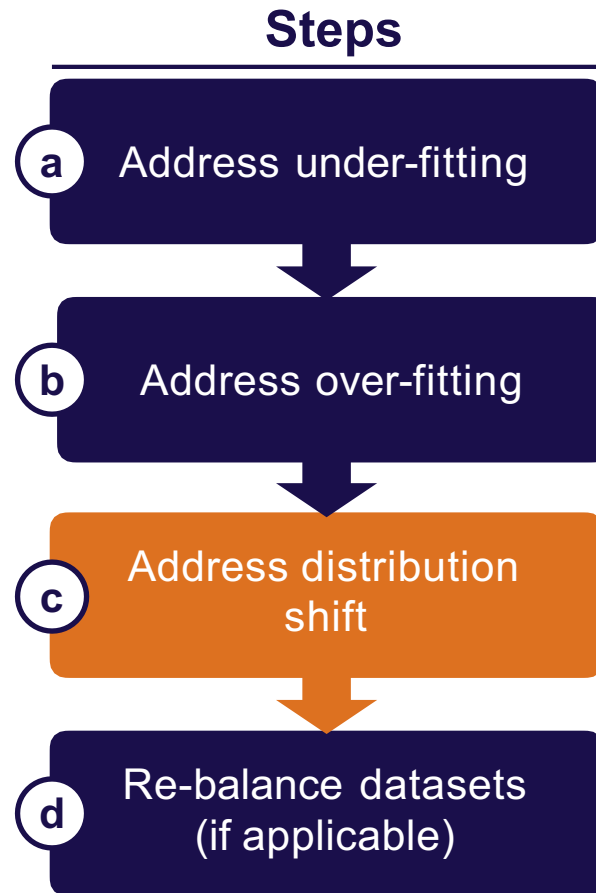


0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy


# Prioritizing improvements (i.e., applied b-v)





# Addressing distribution shift


**Try first**

- 
- A. Analyze test-val set errors & collect more training data to compensate
  - B. Analyze test-val set errors & synthesize more training data to compensate
  - C. Apply domain adaptation techniques to training & test distributions

**Try later**

# Error analysis

## Test-val set errors (no pedestrian detected)





## Train-val set errors (no pedestrian detected)



# Error analysis

## Test-val set errors (no pedestrian detected)




## Train-val set errors (no pedestrian detected)




**Error type 1: hard-to-see pedestrians**

# Error analysis

Test-val set errors (no pedestrian detected)



Train-val set errors (no pedestrian detected)




Error type 2: reflections



# Error analysis

## Test-val set errors (no pedestrian detected)



**Error type 3 (test-val only):**  
night scenes

## Train-val set errors (no pedestrian detected)



# Error analysis

Error type	Error % (train-val)	Error % (test-val)	Potential solutions	Priority
1. Hard-to-see pedestrians	0.1%	0.1%	<ul style="list-style-type: none"><li>• Better sensors</li></ul>	Low
2. Reflections	0.3%	0.3%	<ul style="list-style-type: none"><li>• Collect more data with reflections</li><li>• Add synthetic reflections to train set</li><li>• Try to remove with pre-processing</li><li>• Better sensors</li></ul>	Medium
3. Nighttime scenes	0.1%	1%	<ul style="list-style-type: none"><li>• Collect more data at night</li><li>• Synthetically darken training images</li><li>• Simulate night-time data</li><li>• Use domain adaptation</li></ul>	High

# Domain adaptation

## What is it?

---

Techniques to train on “source” distribution and generalize to another “target” using only unlabeled data or limited labeled data

## When should you consider using it?

---

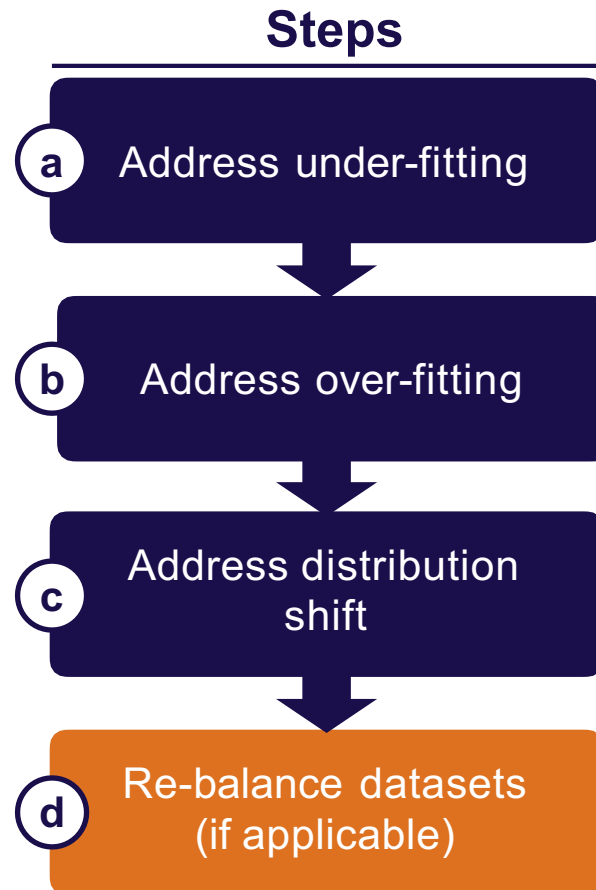
- Access to labeled data from test distribution is limited
- Access to relatively similar data is plentiful

# Types of domain adaptation

Type	Use case	Example techniques
<b>Supervised</b>	You have limited data from target domain	<ul style="list-style-type: none"><li>• Fine-tuning a pre-trained model</li><li>• Adding target data to train set</li></ul>
<b>Un-supervised</b>	You have lots of un-labeled data from target domain	<ul style="list-style-type: none"><li>• Correlation Alignment (CORAL)</li><li>• Domain confusion</li><li>• CycleGAN</li></ul>



# Prioritizing improvements (i.e., applied b-v)

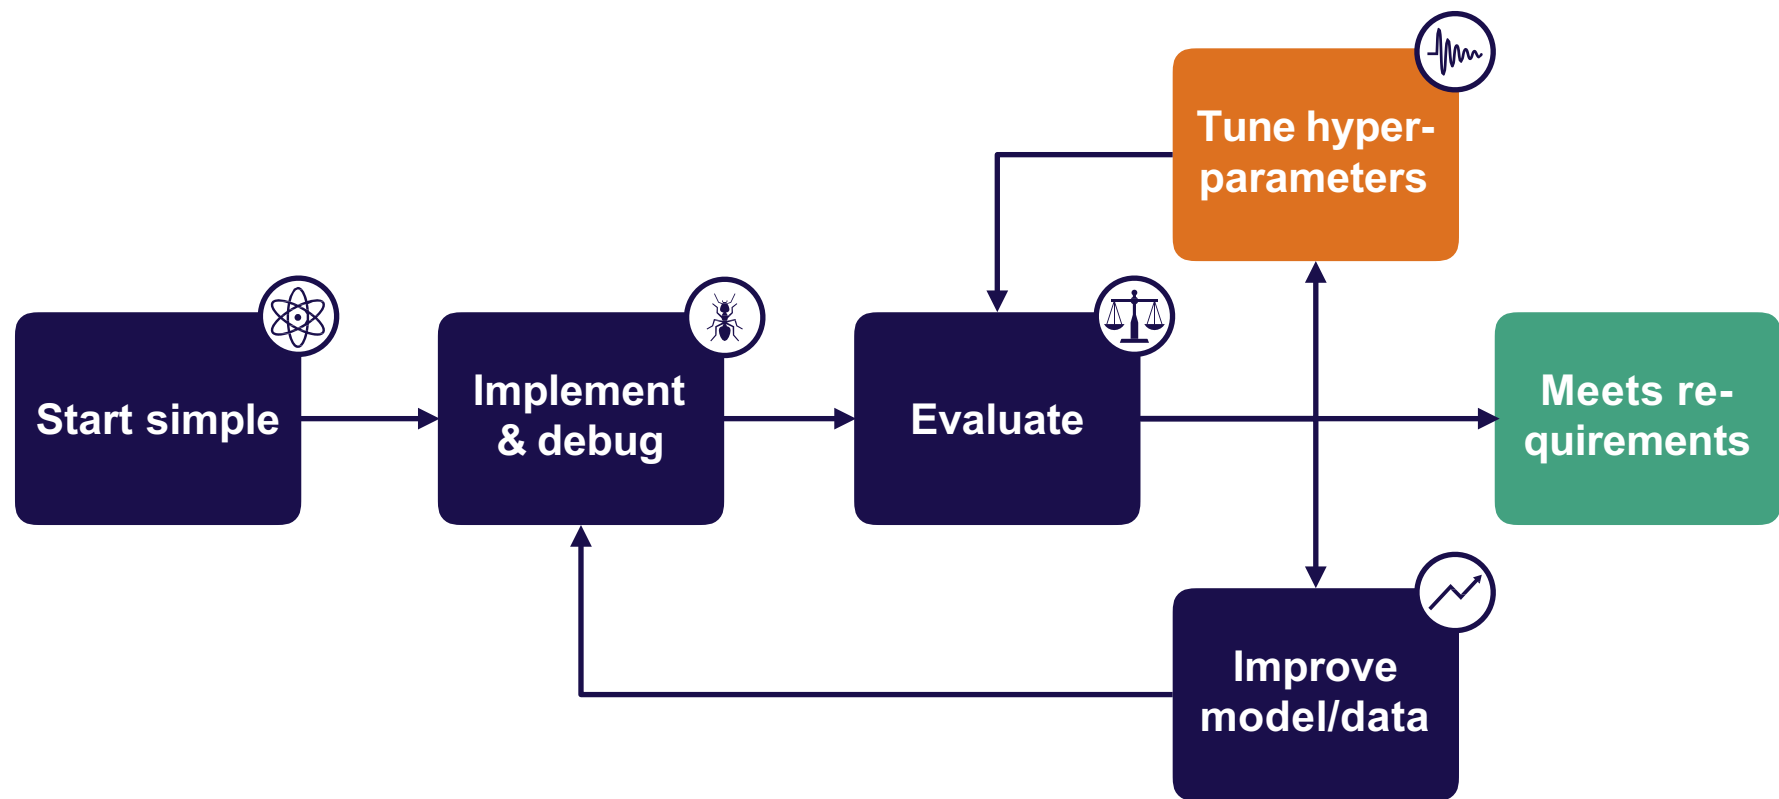


# Rebalancing datasets

- If (test)-val looks significantly better than test, you overfit to the val set
- This happens with small val sets or lots of hyper parameter tuning
- When it does, recollect val data

# Questions?

# Strategy for DL troubleshooting



# Hyperparameter optimization

## Model & optimizer choices?

**Network:** ResNet

- How many layers?
- Weight initialization?
- Kernel size?
- Etc

**Optimizer:** Adam

- Batch size?
- Learning rate?
- beta1, beta2, epsilon?

**Regularization**

- ....

## Running example



0 (no pedestrian)

1 (yes pedestrian)

**Goal:** 99% classification accuracy

# Which hyper-parameters to tune?

## Choosing hyper-parameters

- More sensitive to some than others
- Depends on choice of model
- Rules of thumb (only) to the right
- Sensitivity is relative to default values!  
(e.g., if you are using all-zeros weight initialization or vanilla SGD, changing to the defaults will make a big difference)

Hyperparameter	Approximate sensitivity
Learning rate	High
Learning rate schedule	High
Optimizer choice	Low
Other optimizer params (e.g., Adam beta1)	Low
Batch size	Low
Weight initialization	Medium
Loss function	High
Model depth	Medium
Layer size	High
Layer params (e.g., kernel size)	Medium
Weight of regularization	Medium
Nonlinearity	Low

# Method 1: manual hyperparam optimization

## How it works

---

- Understand the algorithm
  - E.g., higher learning rate means faster less stable training
- Train & evaluate model
- Guess a better hyperparam value & re-evaluate
- Can be combined with other methods (e.g., manually select parameter ranges to optimizer over)

## Advantages

---

- For a skilled practitioner, may require least computation to get good result


## Disadvantages

---

- Requires detailed understanding of the algorithm
- Time-consuming

# Method 2: grid search

## How it works



## Advantages

- Super simple to implement
- Can produce good results


## Disadvantages

- Not very efficient: need to train on all cross-combos of hyperparameters
- May require prior knowledge about parameters to get good results



# Method 3: random search

## How it works



## Advantages


- Easy to implement
- Often produces better results than grid search

## Disadvantages

- Not very interpretable
- May require prior knowledge about parameters to get good results

# Method 4: coarse-to-fine

How it works



---

Advantages


---

Disadvantages

---

# Method 4: coarse-to-fine

How it works



Advantages


---

Disadvantages

---

# Method 4: coarse-to-fine

## How it works



## Advantages


---

## Disadvantages

---

# Method 4: coarse-to-fine

## How it works



## Advantages


---

## Disadvantages

---

# Method 4: coarse-to-fine

## How it works



## Advantages

- Can narrow in on very high performing hyperparameters
- Most used method in practice

## Disadvantages

- Somewhat manual process

# Method 5: Bayesian hyperparam opt

## How it works (at a high level)

---

- Start with a prior estimate of parameter distributions
- Maintain a probabilistic model of the relationship between hyper-parameter values and model performance
- Alternate between:
  - Training with the hyper-parameter values that maximize the expected improvement
  - Using training results to update our probabilistic model
- To learn more, see:

## Advantages

---

- Generally the most efficient hands-off way to choose hyperparameters

## Disadvantages

---

- Difficult to implement from scratch
- Can be hard to integrate with off-the-shelf tools

<https://towardsdatascience.com/a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-b8172278050f>

# Method 5: Bayesian hyperparam opt

## How it works (at a high level)

- Start with a prior estimate of parameter distributions
- Maintain a probabilistic model of the relationship between hyper-parameter values and model performance
- Alternate between:
  - Training with the hyper-parameter values that maximize model performance
  - Using training results to update our probabilistic model
- To learn more, see:

## Advantages

- Generally the most efficient hands-off way to choose hyperparameters

## Disadvantages

- Can be hard to integrate with off-the-shelf tools

More on tools to do this automatically in the infrastructure & tooling lecture!

<https://towardsdatascience.com/a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-b8172278050f>



# Summary of how to optimize hyperparams

- Coarse-to-fine random searches
- Consider Bayesian hyper-parameter optimization solutions as your codebase matures

# Questions?

# Conclusion

- **DL debugging is hard due to many competing sources of error**
- **To train bug-free DL models, we treat building our model as an iterative process**
- **The following steps can make the process easier and catch errors as early as possible**

# How to build bug-free DL models



**Start  
simple**

- Choose the simplest model & data possible (e.g., LeNet on a subset of your data)



**Implement  
& debug**

- Once model runs, overfit a single batch & reproduce a known result



**Evaluate**

- Apply the bias-variance decomposition to decide what to do next



**Tune hyp-  
eparams**

- Use coarse-to-fine random searches



**Improve  
model/data**

- Make your model bigger if you underfit; add data or regularize if you overfit

Thank you!